



Choupo

The Open Source Chemical Process Simulator

Process Design Heuristics

Choupo-dev

Copyright © 2026 Vítor Geraldès.

Authors: Vítor Geraldès.

This work is licensed under a Creative Commons Attribution-ShareAlike 4.0 International License.

Code examples, Choupo case files, and other machine-readable examples are licensed under GPL-3.0-or-later.

Typeset in L^AT_EX.

Acknowledgment. Choupo was conceived and architected by Vítor Geraldès. Substantial parts of the initial implementation, documentation, and tutorial corpus were produced with assistance from Anthropic Claude Code using Claude Opus/Fable models. This guide is published under human curation, review, and editorial responsibility by Vítor Geraldès.

License

Attribution-ShareAlike 4.0 International

=====

Creative Commons Corporation ("Creative Commons") is not a law firm and does not provide legal services or legal advice. Distribution of Creative Commons public licenses does not create a lawyer-client or other relationship. Creative Commons makes its licenses and related information available on an "as-is" basis. Creative Commons gives no warranties regarding its licenses, any material licensed under their terms and conditions, or any related information. Creative Commons disclaims all liability for damages resulting from their use to the fullest extent possible.

Using Creative Commons Public Licenses

Creative Commons public licenses provide a standard set of terms and conditions that creators and other rights holders may use to share original works of authorship and other material subject to copyright and certain other rights specified in the public license below. The following considerations are for informational purposes only, are not exhaustive, and do not form part of our licenses.

Considerations for licensors: Our public licenses are intended for use by those authorized to give the public permission to use material in ways otherwise restricted by copyright and certain other rights. Our licenses are irrevocable. Licensors should read and understand the terms and conditions of the license they choose before applying it. Licensors should also secure all rights necessary before applying our licenses so that the public can reuse the material as expected. Licensors should clearly mark any material not subject to the license. This includes other CC-licensed material, or material used under an exception or limitation to copyright. More considerations for licensors: wiki.creativecommons.org/Considerations_for_licensors

Considerations for the public: By using one of our public licenses, a licensor grants the public permission to use the licensed material under specified terms and conditions. If the licensor's permission is not necessary for any reason--for example, because of any applicable exception or limitation to copyright--then that use is not regulated by the license. Our licenses grant only permissions under copyright and certain other rights that a licensor has authority to grant. Use of the licensed material may still be restricted for other reasons, including because others have copyright or other rights in the material. A licensor may make special requests, such as asking that all changes be marked or described. Although not required by our licenses, you are encouraged to respect those requests where reasonable. More considerations for the public: wiki.creativecommons.org/Considerations_for_licensees

=====

Creative Commons Attribution-ShareAlike 4.0 International Public License

By exercising the Licensed Rights (defined below), You accept and agree to be bound by the terms and conditions of this Creative Commons Attribution-ShareAlike 4.0 International Public License ("Public License"). To the extent this Public License may be interpreted as a contract, You are granted the Licensed Rights in consideration of Your acceptance of these terms and conditions, and the Licensor grants You such rights in consideration of benefits the Licensor receives from making the Licensed Material available under these terms and conditions.

Section 1 -- Definitions.

- a. Adapted Material means material subject to Copyright and Similar Rights that is derived from or based upon the Licensed Material and in which the Licensed Material is translated, altered, arranged, transformed, or otherwise modified in a manner requiring permission under the Copyright and Similar Rights held by the Licensor. For purposes of this Public License, where the Licensed Material is a musical work, performance, or sound recording, Adapted Material is always produced where the Licensed Material is synched in timed relation with a moving image.
- b. Adapter's License means the license You apply to Your Copyright and Similar Rights in Your contributions to Adapted Material in accordance with the terms and conditions of this Public License.
- c. BY-SA Compatible License means a license listed at creativecommons.org/compatiblelicenses, approved by Creative Commons as essentially the equivalent of this Public License.
- d. Copyright and Similar Rights means copyright and/or similar rights closely related to copyright including, without limitation, performance, broadcast, sound recording, and Sui Generis Database Rights, without regard to how the rights are labeled or categorized. For purposes of this Public License, the rights specified in Section 2(b)(1)-(2) are not Copyright and Similar

Rights.

- e. Effective Technological Measures means those measures that, in the absence of proper authority, may not be circumvented under laws fulfilling obligations under Article 11 of the WIPO Copyright Treaty adopted on December 20, 1996, and/or similar international agreements.
- f. Exceptions and Limitations means fair use, fair dealing, and/or any other exception or limitation to Copyright and Similar Rights that applies to Your use of the Licensed Material.
- g. License Elements means the license attributes listed in the name of a Creative Commons Public License. The License Elements of this Public License are Attribution and ShareAlike.
- h. Licensed Material means the artistic or literary work, database, or other material to which the Licensor applied this Public License.
- i. Licensed Rights means the rights granted to You subject to the terms and conditions of this Public License, which are limited to all Copyright and Similar Rights that apply to Your use of the Licensed Material and that the Licensor has authority to license.
- j. Licensor means the individual(s) or entity(ies) granting rights under this Public License.
- k. Share means to provide material to the public by any means or process that requires permission under the Licensed Rights, such as reproduction, public display, public performance, distribution, dissemination, communication, or importation, and to make material available to the public including in ways that members of the public may access the material from a place and at a time individually chosen by them.
- l. Sui Generis Database Rights means rights other than copyright resulting from Directive 96/9/EC of the European Parliament and of the Council of 11 March 1996 on the legal protection of databases, as amended and/or succeeded, as well as other essentially equivalent rights anywhere in the world.
- m. You means the individual or entity exercising the Licensed Rights under this Public License. Your has a corresponding meaning.

Section 2 -- Scope.

- a. License grant.
 - 1. Subject to the terms and conditions of this Public License, the Licensor hereby grants You a worldwide, royalty-free, non-sublicensable, non-exclusive, irrevocable license to exercise the Licensed Rights in the Licensed Material to:
 - a. reproduce and Share the Licensed Material, in whole or in part; and
 - b. produce, reproduce, and Share Adapted Material.
 - 2. Exceptions and Limitations. For the avoidance of doubt, where Exceptions and Limitations apply to Your use, this Public License does not apply, and You do not need to comply with its terms and conditions.
 - 3. Term. The term of this Public License is specified in Section 6(a).
 - 4. Media and formats; technical modifications allowed. The Licensor authorizes You to exercise the Licensed Rights in all media and formats whether now known or hereafter created, and to make technical modifications necessary to do so. The Licensor waives and/or agrees not to assert any right or authority to forbid You from making technical modifications necessary to exercise the Licensed Rights, including technical modifications necessary to circumvent Effective Technological Measures. For purposes of this Public License, simply making modifications authorized by this Section 2(a) (4) never produces Adapted Material.
 - 5. Downstream recipients.
 - a. Offer from the Licensor -- Licensed Material. Every recipient of the Licensed Material automatically receives an offer from the Licensor to exercise the Licensed Rights under the terms and conditions of this Public License.
 - b. Additional offer from the Licensor -- Adapted Material. Every recipient of Adapted Material from You automatically receives an offer from the Licensor to exercise the Licensed Rights in the Adapted Material under the conditions of the Adapter's License You apply.
 - c. No downstream restrictions. You may not offer or impose any additional or different terms or conditions on, or apply any Effective Technological Measures to, the Licensed Material if doing so restricts exercise of the Licensed Rights by any recipient of the Licensed Material.

6. No endorsement. Nothing in this Public License constitutes or may be construed as permission to assert or imply that You are, or that Your use of the Licensed Material is, connected with, or sponsored, endorsed, or granted official status by, the Licensor or others designated to receive attribution as provided in Section 3(a)(1)(A)(i).

b. Other rights.

1. Moral rights, such as the right of integrity, are not licensed under this Public License, nor are publicity, privacy, and/or other similar personality rights; however, to the extent possible, the Licensor waives and/or agrees not to assert any such rights held by the Licensor to the limited extent necessary to allow You to exercise the Licensed Rights, but not otherwise.
2. Patent and trademark rights are not licensed under this Public License.
3. To the extent possible, the Licensor waives any right to collect royalties from You for the exercise of the Licensed Rights, whether directly or through a collecting society under any voluntary or waivable statutory or compulsory licensing scheme. In all other cases the Licensor expressly reserves any right to collect such royalties.

Section 3 -- License Conditions.

Your exercise of the Licensed Rights is expressly made subject to the following conditions.

a. Attribution.

1. If You Share the Licensed Material (including in modified form), You must:
 - a. retain the following if it is supplied by the Licensor with the Licensed Material:
 - i. identification of the creator(s) of the Licensed Material and any others designated to receive attribution, in any reasonable manner requested by the Licensor (including by pseudonym if designated);
 - ii. a copyright notice;
 - iii. a notice that refers to this Public License;
 - iv. a notice that refers to the disclaimer of warranties;
 - v. a URI or hyperlink to the Licensed Material to the extent reasonably practicable;
 - b. indicate if You modified the Licensed Material and retain an indication of any previous modifications; and
 - c. indicate the Licensed Material is licensed under this Public License, and include the text of, or the URI or hyperlink to, this Public License.
2. You may satisfy the conditions in Section 3(a)(1) in any reasonable manner based on the medium, means, and context in which You Share the Licensed Material. For example, it may be reasonable to satisfy the conditions by providing a URI or hyperlink to a resource that includes the required information.
3. If requested by the Licensor, You must remove any of the information required by Section 3(a)(1)(A) to the extent reasonably practicable.

b. ShareAlike.

In addition to the conditions in Section 3(a), if You Share Adapted Material You produce, the following conditions also apply.

1. The Adapter's License You apply must be a Creative Commons license with the same License Elements, this version or later, or a BY-SA Compatible License.
2. You must include the text of, or the URI or hyperlink to, the Adapter's License You apply. You may satisfy this condition in any reasonable manner based on the medium, means, and context in which You Share Adapted Material.
3. You may not offer or impose any additional or different terms or conditions on, or apply any Effective Technological Measures to, Adapted Material that restrict exercise of the rights granted under the Adapter's License You apply.

Section 4 -- Sui Generis Database Rights.

Where the Licensed Rights include Sui Generis Database Rights that apply to Your use of the Licensed Material:

- a. for the avoidance of doubt, Section 2(a)(1) grants You the right to extract, reuse, reproduce, and Share all or a substantial portion of the contents of the database;
- b. if You include all or a substantial portion of the database contents in a database in which You have Sui Generis Database Rights, then the database in which You have Sui Generis Database Rights (but not its individual contents) is Adapted Material, including for purposes of Section 3(b); and
- c. You must comply with the conditions in Section 3(a) if You Share all or a substantial portion of the contents of the database.

For the avoidance of doubt, this Section 4 supplements and does not replace Your obligations under this Public License where the Licensed Rights include other Copyright and Similar Rights.

Section 5 -- Disclaimer of Warranties and Limitation of Liability.

- a. UNLESS OTHERWISE SEPARATELY UNDERTAKEN BY THE LICENSOR, TO THE EXTENT POSSIBLE, THE LICENSOR OFFERS THE LICENSED MATERIAL AS-IS AND AS-AVAILABLE, AND MAKES NO REPRESENTATIONS OR WARRANTIES OF ANY KIND CONCERNING THE LICENSED MATERIAL, WHETHER EXPRESS, IMPLIED, STATUTORY, OR OTHER. THIS INCLUDES, WITHOUT LIMITATION, WARRANTIES OF TITLE, MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE, NON-INFRINGEMENT, ABSENCE OF LATENT OR OTHER DEFECTS, ACCURACY, OR THE PRESENCE OR ABSENCE OF ERRORS, WHETHER OR NOT KNOWN OR DISCOVERABLE. WHERE DISCLAIMERS OF WARRANTIES ARE NOT ALLOWED IN FULL OR IN PART, THIS DISCLAIMER MAY NOT APPLY TO YOU.
- b. TO THE EXTENT POSSIBLE, IN NO EVENT WILL THE LICENSOR BE LIABLE TO YOU ON ANY LEGAL THEORY (INCLUDING, WITHOUT LIMITATION, NEGLIGENCE) OR OTHERWISE FOR ANY DIRECT, SPECIAL, INDIRECT, INCIDENTAL, CONSEQUENTIAL, PUNITIVE, EXEMPLARY, OR OTHER LOSSES, COSTS, EXPENSES, OR DAMAGES ARISING OUT OF THIS PUBLIC LICENSE OR USE OF THE LICENSED MATERIAL, EVEN IF THE LICENSOR HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH LOSSES, COSTS, EXPENSES, OR DAMAGES. WHERE A LIMITATION OF LIABILITY IS NOT ALLOWED IN FULL OR IN PART, THIS LIMITATION MAY NOT APPLY TO YOU.
- c. The disclaimer of warranties and limitation of liability provided above shall be interpreted in a manner that, to the extent possible, most closely approximates an absolute disclaimer and waiver of all liability.

Section 6 -- Term and Termination.

- a. This Public License applies for the term of the Copyright and Similar Rights licensed here. However, if You fail to comply with this Public License, then Your rights under this Public License terminate automatically.
- b. Where Your right to use the Licensed Material has terminated under Section 6(a), it reinstates:
 - 1. automatically as of the date the violation is cured, provided it is cured within 30 days of Your discovery of the violation; or
 - 2. upon express reinstatement by the Licensor.

For the avoidance of doubt, this Section 6(b) does not affect any right the Licensor may have to seek remedies for Your violations of this Public License.

- c. For the avoidance of doubt, the Licensor may also offer the Licensed Material under separate terms or conditions or stop distributing the Licensed Material at any time; however, doing so will not terminate this Public License.
- d. Sections 1, 5, 6, 7, and 8 survive termination of this Public License.

Section 7 -- Other Terms and Conditions.

- a. The Licensor shall not be bound by any additional or different terms or conditions communicated by You unless expressly agreed.
- b. Any arrangements, understandings, or agreements regarding the Licensed Material not stated herein are separate from and independent of the terms and conditions of this Public License.

Section 8 -- Interpretation.

- a. For the avoidance of doubt, this Public License does not, and shall not be interpreted to, reduce, limit, restrict, or impose conditions on any use of the Licensed Material that could lawfully be made without permission under this Public License.
- b. To the extent possible, if any provision of this Public License is deemed unenforceable, it shall be automatically reformed to the minimum extent necessary to make it enforceable. If the provision cannot be reformed, it shall be severed from this Public License without affecting the enforceability of the remaining terms and conditions.

- c. No term or condition of this Public License will be waived and no failure to comply consented to unless expressly agreed to by the Licensor.
- d. Nothing in this Public License constitutes or may be interpreted as a limitation upon, or waiver of, any privileges and immunities that apply to the Licensor or You, including from the legal processes of any jurisdiction or authority.

=====
Creative Commons is not a party to its public licenses. Notwithstanding, Creative Commons may elect to apply one of its public licenses to material it publishes and in those instances will be considered the "Licensor." The text of the Creative Commons public licenses is dedicated to the public domain under the CC0 Public Domain Dedication. Except for the limited purpose of indicating that material is shared under a Creative Commons public license or as otherwise permitted by the Creative Commons policies published at creativecommons.org/policies, Creative Commons does not authorize the use of the trademark "Creative Commons" or any other trademark or logo of Creative Commons without its prior written consent including, without limitation, in connection with any unauthorized modifications to any of its public licenses or any other arrangements, understandings, or agreements concerning use of licensed material. For the avoidance of doubt, this paragraph does not form part of the public licenses.

Creative Commons may be contacted at creativecommons.org.

Trademarks

Choupo is a trademark of TalentGround Lda.

Anthropic, Claude, and Claude Code are trademarks or registered trademarks of Anthropic, PBC.

Disclaimer & scope — read before using any number in this guide

These are **engineering rules of thumb** — deliberately approximate, order-of-magnitude starting points for conceptual design and teaching, stated as ranges. They are **not** a substitute for, and must never override:

- the **governing codes and standards** (ASME, API, NFPA, IEC, ISO, EN, and local regulations);
- **detailed design and review by a qualified, licensed professional engineer**;
- a formal **process-hazard analysis** (HAZOP / LOPA) and the **manufacturer's data** for the specific equipment.

Nothing here is a recommendation for a real plant, and no value has been validated for any specific service. **Safety-critical decisions** — pressure relief, fired equipment, toxic or flammable inventories, electrical area classification, effluent limits — *can cost lives or breach the law if wrong* and **MUST** follow the applicable standard and competent review, never a rule of thumb. Where a number could touch safety or the environment, the controlling standard is named (see *Standards & further reading*).

To the maximum extent permitted by law, **Choupo, its authors, and its contributors accept no liability** for any use of this material; it is provided “as is”, without warranty of any kind, for education. Professional responsibility for any design remains entirely with the engineer of record.

Contents

1	Process synthesis & flowsheet structure	2
2	Reactors & reaction engineering	3
2.1	Type selection	3
2.2	How far to push conversion	3
2.3	Temperature, pressure & exotherm policy	3
2.4	Sizing rules	3
3	Distillation & separation-train design	4
3.1	Internals selection (trays vs packing)	4
3.2	Diameter from vapour load (flooding)	4
3.3	Reflux and FUG sizing	4
3.4	Sequencing multicomponent splits	5
3.5	Azeotropes	5
4	Absorption, stripping & gas treating	6
5	Liquid-liquid extraction	7
6	Membranes & aqueous/advanced separations	8
6.1	Selecting the membrane class	8
6.2	Recovery vs scaling — the master trade-off	8
6.3	Crossflow, flux and polarisation	8
6.4	Staging, trains, pretreatment	8
7	Solids processing & particle technology	9
7.1	Crystallisation route — pick by the solubility curve	9
7.2	Solid–liquid separation — filtration vs centrifugation	9
7.3	Dryer selection — by feed state and particle size	9
7.4	Gas–solid separation & fluidisation velocities	9
8	Evaporation & multiple-effect	10
9	Gas-liquid separators & process drums	11
10	Heat exchangers — type selection & sizing	12
10.1	Type-selection table	12
10.2	Gas-phase service	12
10.3	Allocation & velocities	12
10.4	Phase-change service	12
10.5	$\Delta T_{\min}$, LMTD correction, fouling	13
11	Heat integration & pinch analysis	14
12	Fired heaters & furnaces	15
13	Compressors, expanders & pressure change	17
13.1	TYPE SELECTION	17
13.2	Stages, ratio per stage & intercooling	17
13.3	Efficiency, surge & turndown	17
13.4	When an expander pays back	17
14	Drivers: electric motors, steam & gas turbines	18
15	Fluid transport, line sizing & hydraulics	19
15.1	Driver selection by pressure ratio	19
15.2	Recommended superficial velocities	19
15.3	Economic pipe diameter & ΔP budget	19
15.4	NPSH (the pump killer)	19
16	Control valves, pressure letdown & choking	20

17 Refrigeration & vacuum systems	21
17.1 Refrigeration — pick the warmest level that works	21
17.2 Vacuum — match the pump to the absolute pressure	21
18 Storage & intermediate tanks	22
19 Process control & instrumentation heuristics	23
20 Pressure relief, flare & safety basics	25
21 Effluent treatment — liquid & gaseous	27
21.1 Liquid effluent (wastewater)	27
21.2 Gaseous effluent (air emissions)	27
22 Cost estimation & profitability	28
23 Materials of construction & mechanical design	30
24 Standards & further reading	32
25 Equipment selection at a glance	33
26 Equipment selection at a glance	34

This guide is Choupo's answer to the rules-of-thumb appendices you reach for in Turton, Seider, and Sinnott & Towler — but rebuilt for a simulator whose whole point is that *nothing is hidden*. Where a textbook table hands you a single number ("use $\Delta T_{\min} = 10\text{ }^{\circ}\text{C}$ ", "size liquid lines at 1–3 m/s"), this guide hands you a **range, a why, and a runnable hypothesis**. The eight faculty sections that follow cover the full sweep of decisions a student faces when a problem stops being "solve this flash" and becomes "what equipment do I even pick?": process structure, reactor type, separation train, exchanger type and velocity, line sizing and ΔP budgets, membranes, solids handling, and the metal envelope that contains all of it.

How to read every rule here: a heuristic is a *hypothesis*, not a verdict. "Pick a PFR when you need high conversion of a single reaction" (the reactor rule below) is a starting bet that buys you 80 % of the answer in 20 % of the effort — it is *not* a substitute for the run. Choupo is glass-box precisely so you can test the bet: build the case, watch the Newton iterations, read the KPIs, and confirm (or refute) the rule for *your* feed, *your* thermodynamics, *your* pressure level. Wherever a rule maps cleanly onto an existing tutorial, the section points you at it — run it, perturb it, and see the range move. A heuristic you have personally falsified once is worth ten you merely memorised.

Why the sections are ordered the way they are — the design onion. Process design proceeds from the inside out: the reactor sits at the core, the separation-and-recycle system wraps it, the heat-recovery network wraps that, and the utilities, hydraulics, and mechanical envelope wrap the whole. Each layer's decisions *constrain* the next — the reactor's conversion fixes what the separation train must handle; the separation train's streams fix the heating and cooling duties; those duties fix the pumping power and line sizes; and the worst temperature, pressure, and corrosivity anywhere inside fix the materials of construction on the outside. So we read **outermost framing first** (synthesis), then **inward to the core** (reactors), then **back out through the layers**: the full separation train kept together (distillation → membranes → solids), then heat recovery (heat exchangers), then the connective hydraulics, and finally the material envelope that contains everything. This keeps every decision in the order in which its inputs actually become available.

1 Process synthesis & flowsheet structure

Design the flowsheet **inside-out**, following Douglas' onion. Each layer fixes the degrees of freedom (DOF) of the next, so never optimise an inner layer against an outer-layer guess.

Onion layer	What you DECIDE	Heuristic numbers / ranges	Physical/economic WHY
1. Reactor	CSTR vs PFR, single-pass conversion X	Push X to 0.6–0.9 for cheap reactant; stop where recycle cost rises faster than reactor cost. Keep X ≤ 0.95 if selectivity falls (consecutive rxns)	Reactor cost $\propto V \propto \ln[1/(1-X)]$ for a first-order PFR; recycle/separation cost $\propto$ unconverted feed. The minimum-TAC X sits where $\partial(\text{reactor})/\partial X = -\partial(\text{recycle})/\partial X$
2. Separation	Recover vs recycle; sequence	Recover species worth $> \sim 3\text{--}5\times$ its separation cost; aim 99%+ recovery of valuable product	Lost product is pure margin; over-purifying a recycle wastes reboiler duty
3. Recycle/purge	Purge fraction for inerts	Purge so inert mole-fraction in loop $\leq$ 0.10–0.30 ; purge loss of reactant $< 1\text{--}5\%$ of feed	Inert build-up raises recycle compressor/pump duty ($\propto$ total recycle flow); too-small purge $\rightarrow$ runaway loop inventory
4. Heat recovery	HEN, $\Delta T_{\min}$	$\Delta T_{\min}$ 10 °C liquid–liquid, 20–40 °C with a gas stream, 3–5 °C cryogenic/refrigerated	Area $\propto 1/\Delta T_{\min}$ ($U_{\text{gas}} \sim 15\text{--}60 \text{ W/m}^2\text{K}$ vs $U_{\text{liq}} \sim 300\text{--}1000$); small ΔT buys energy but explodes area cost
5. Utilities	Hot/cold service, fired vs salt	LP/MP/HP steam below $\sim 250 \text{ °C}$; Dowtherm A to $\sim 400 \text{ °C}$; HiTec salt 142–535 °C ; fired heater only above salt range	Match utility T just above pinch; over-hot utility destroys exergy and over-sizes ΔT

Input–output structure first. Before any unit, draw the I/O box: feeds in, products + purge + by-products out. Set **feed stoichiometric ratio** (slight excess, $\sim 1.05\text{--}1.2$, of the cheaper/limiting-safety reactant), decide which species **leave vs recycle**, and locate the **purge** (needed whenever a feed impurity or light by-product has no other exit). This fixes overall material balance DOF *before* you size anything.

Reactor type. Use a **PFR** when you need high conversion of a single reaction (plug flow $\rightarrow$ smaller volume for given X) or want a temperature/composition profile; use a **CSTR** for liquid-phase, exothermic, fast, or solids/biological systems needing isothermality and mixing. For series-parallel selectivity, a **CSTR-then-PFR** or staged CSTRs (3–4 stages approach PFR) is common. See **tutorials/steady/reactors** cases; for economics-driven X, `optim05_reactor_npv`.

Distillation sequencing (Lightest-out / direct sequence is the default): (1) remove **most plentiful** component early; (2) take the **lightest** (highest relative volatility, α) split first when $\alpha > \sim 1.5$; (3) leave the **hardest split** ($\alpha \rightarrow 1$) and the **highest-purity** product for last, alone, to avoid re-contaminating it; (4) favour **near-50/50** splits to balance column loads. Each ordering is a hypothesis — run the Wang-Henke / FUG shortcut cases (**tutorials/steady/distillation**) and compare reboiler+condenser duty.

Degrees of freedom discipline. A recycle adds one tear stream; Choupo's Newton-on-tears solver converges it, but **you** still own the design specs (purge fraction, reactor X, reflux). Count them: overall DOF = stream variables — independent balances — fixed specs. If the flowsheet won't close, you've under- or over-specified — the glass-box log tells you which. Sequence the whole sheet **reactor** $\rightarrow$ **separation** $\rightarrow$ **recycle** $\rightarrow$ **HEN** $\rightarrow$ **utilities**, re-pinching only after the recycle converges, because recycle flows change the composite curves.

2 Reactors & reaction engineering

The reactor is rarely the most expensive vessel, but it sets the composition every downstream unit must clean up — so its selection ripples through the whole flowsheet cost. Pick the **type** from mixing/kinetics, then size from **residence time** and **heat load**.

2.1 Type selection

Criterion	Batch / semi-batch	CSTR	PFR / tubular	Fixed bed (catalytic)	Fluidised bed
Throughput	< 500 t/yr, multiproduct	medium, liquid	medium–high	high, gas+solid cat.	high, fast- deactivating cat.
Phase	liquid, slurry	liquid, L–L, gas–liq	gas or liquid	gas/liquid + solid	gas + fine solid
Kinetics favoured	slow (hrs)	any; needs CSTR robustness	high order, fast	heterogeneous	very exo/en- dothermic, ΔT -sensitive
Conversion per pass	high (time-driven)	lowest for given V	high	high	moderate
Heat control	jacket/coil, batchwise	excellent (isothermal, well-mixed)	hard (axial profile)	hot-spots; multitube	near-isothermal (solids mix)
τ for same X	—	largest	smallest	small	small

The governing trade-off: a CSTR operates at the **outlet** (low-reactant) concentration, so for $n > 0$ kinetics it needs the largest volume — typically **1.5–4×** a PFR for the same conversion (Choupo `cstr01_first_order` vs `pfr01_first_order` shows this head-to-head). You accept that penalty to buy near-perfect temperature control and the ability to run slurries/fouling systems a tube would plug. Use **CSTRs-in-series** (3–5 tanks) to approach plug-flow behaviour while keeping the mixing robustness.

2.2 How far to push conversion

Per-pass conversion is an **optimisation, not a maximisation**. Reactor volume scales roughly as $V \propto \ln[1/(1-X)]$ (first order) — the last few percent of conversion cost disproportionate volume and, worse, residence time that breeds side products and erodes selectivity. **Heuristics:**

- Single-pass irreversible, valuable feed: target **X = 0.6–0.9**, then **recycle unconverted reactant** (recovery + recycle is almost always cheaper than the extra reactor volume; see `conversion01_hda` and the recycle flowsheet `process01_reactor_flash`).
- Reversible reactions: stop at **80–95% of equilibrium X**; chasing equilibrium needs infinite τ (`cstr02_reversible_wgs`, `pfr02_reversible_wgs`).
- The economic optimum trades reactor capital vs separation/recycle duty — run a sweep; `optim05_reactor_npv` does exactly this on NPV.

2.3 Temperature, pressure & exotherm policy

Run as **hot as selectivity and catalyst/material limits allow** — rate roughly doubles per **+10 °C** — but watch the Arrhenius gap: if the side reaction has higher E_a , hotter means *less* selective. For reversible exothermic reactions use a **declining T profile** (hot inlet for rate, cool outlet for equilibrium).

Heat removal sets the geometry. Estimate **adiabatic $\Delta T_{ad} = (-\Delta H_{rxn}) \cdot X \cdot C_{feed} / (\rho \cdot cp)$** ; if ΔT_{ad} exceeds **~20–50 °C**, an adiabatic vessel is unsafe — go to cooled multitube, CSTR with coil/jacket, or dilute/recycle. Jacket U for stirred liquids $\approx$ **300–1000 W/m²K**; internal coils **500–1500**. Keep the **adiabatic-temperature-rise / runaway** margin honest: `batch06_adaptive_runaway` demonstrates thermal runaway when cooling can't keep up.

2.4 Sizing rules

- **Liquid CSTR:** τ from kinetics; vessel **H/D $\approx$ 1–3**, fill to **70–85%**, freeboard for disengagement. Confirm against `batch01_first_order`.
- **Catalytic fixed bed:** superficial gas velocity **0.3–1.5 m/s**; keep **$\Delta P \leq 10$ –20%** of inlet P (Ergun); bed **L/D ≥ 5** , particle **3–6 mm** to balance ΔP vs internal diffusion.
- **Fluidised bed:** operate at **2–7×** **U_{mf}** , typically **0.3–0.7 m/s**; tall freeboard to cut entrainment.
- **Space velocity** sanity check: gas catalytic **GHSV 500–10,000 h^{−1}**; liquid **LHSV 0.5–5 h^{−1}**.

3 Distillation & separation-train design

Distillation is the workhorse: ~90% of liquid separations, but also the plant's largest single energy sink (reboiler duty). Every heuristic below is a **range to verify** — run the case, read the diameter and reflux the solver reports, and decide.

3.1 Internals selection (trays vs packing)

The *why*: trays handle fouling, solids and wide turndown; structured packing gives low pressure drop and high efficiency per metre but is intolerant of fouling and expensive.

Service / criterion	Pick	Typical numbers	Why
D > 0.8 m, clean, normal P	Sieve/valve trays	spacing 0.45–0.6 m, $\Delta P \approx 0.7$ kPa/tray	cheap, robust, wide turndown (valve 4–5:1)
D < 0.8 m	Random packing	HETP $\approx$ column D	trays uneconomic to install in small shells
Vacuum (P < 0.3 bar)	Structured packing	ΔP 0.1–0.4 kPa/m vs 0.7 kPa/tray	low ΔP keeps reboiler T (and thermal-degradation) down
High purity / many stages	Structured packing	HETP 0.3–0.5 m, a 250–350 m ² /m ³	most stages per metre of height
Fouling / solids / corrosive	Trays	—	cleanable, replaceable
Foaming, low ΔP	Random (Pall/IMTP)	HETP 0.5–0.9 m	open area tolerates surging

3.2 Diameter from vapour load (flooding)

Size on the controlling (usually top) vapour rate. Use the **C-factor** / **Souders-Brown** velocity: $u_{\text{flood}} = C_{\text{sb}} \cdot \sqrt{((\rho_L - \rho_V)/\rho_V)}$, $C_{\text{sb}} \approx 0.07\text{--}0.12$ m/s at 0.45–0.6 m spacing. **Design at 75–85% of flood** (70% for foaming systems). Equivalently the **F-factor** $F = u \cdot \sqrt{\rho_V}$ should land at **1.2–2.5 (m/s)(kg/m³)^{0.5}** for trays, up to ~2.5–3 for structured packing. Typical superficial vapour velocities: **0.5–1.5 m/s** at atmospheric, **2–4 m/s** under vacuum (low ρ_V). Then $A = V/u$, $D = \sqrt{(4A/\pi)}$. Aspect ratio $H/D \approx 8\text{--}30$; if $H/D > 30$ split into two shells. Weir loading and downcomer set the lower bound — turndown below ~50% of flood weeps. `column01_benzene_toluene` reports diameter and the flood fraction; halve the feed and watch it shrink.

Why the top and the bottom of a column want different diameters. The flooding diameter follows the **vapour VOLUMETRIC flow**, not the molar one, and the two sections differ in every term that turns one into the other: temperature, pressure, vapour molar mass and usually the vapour molar flow itself. Below the feed the liquid traffic is larger, which raises the flow parameter $F_{\text{LV}} = (L/V)\sqrt{(\rho_V/\rho_L)}$ and therefore **LOWERS** Fair's capacity parameter C_{sb} — so a stripping tray floods at a lower velocity than a rectifying one. Whether that makes the bottom the wider section is not a rule: the vapour volumetric flow may fall faster than the capacity does, and then the top governs. It is decided by whichever tray loses the argument, tray by tray, which is exactly what `column09_tray_hydraulics` prints. It reports **1.21 m** for its rectifying section against **1.32 m** for its stripping section, on the same trays and the same spacing.

Swaging is an ECONOMIC decision, not a thermodynamic one. Having two diameters, you may build a **swaged** tower — a narrow section, a transition cone and a wide section — or one **straight** shell at the larger diameter. The straight shell buys plate you do not need; the swaged one buys a transition, its fabrication, its extra weld inspection and a second set of tray sizes to hold in stores. The usual rule of thumb: **swage only when the two sections differ by roughly 15–20% in diameter**, and build straight below that, because under about 15% the transition costs more than the plate it saves. Judge it on **DIAMETER**, not on area — 20% in diameter is already 44% in cross-section. The 8% gap in the case above is comfortably inside the band, so that tower is built straight at 1.32 m, and Choupo says so without deciding for you: it publishes both section diameters and their difference and prices no transition, because it has no correlation for one.

3.3 Reflux and FUG sizing

Shortcut design (Fenske-Underwood-Gilliland, `shortcut01_benzene_toluene`): Fenske gives N_{min} at total reflux, Underwood gives R_{min} , Gilliland interpolates. **Heuristic $R = 1.1\text{--}1.5 \cdot R_{\text{min}}$** — the cost optimum, because below 1.1 N (and column height) blows up, above 1.5 reboiler/condenser duty and diameter grow with little stage saving. $N \approx 2 \cdot N_{\text{min}}$ at that reflux. Feed near the stage matching its composition. Validate the shortcut against the rigorous MESH run (`column02_simultaneous`).

3.4 Sequencing multicomponent splits

Direct sequence (lights overhead first) dominates. Heuristics, in priority: (1) remove **corrosive/hazardous** components early; (2) do the **most-plentiful** product split early; (3) save the **hardest split** ($\alpha < 1.2$) and the **highest-purity** split for last, alone; (4) favour **near-50/50** splits; (5) "lightest-first" when relative volatilities are ordered. Each saves reboiler duty by not re-boiling already-separated material.

3.5 Azeotropes

When $\alpha \rightarrow 1$ or an azeotrope blocks the split, ordinary distillation cannot cross it. Wang-Henke (bubble-point) goes **unstable near azeotropes** — **switch to method simultaneous**; (rigorous MESH Newton, `column03_azeotrope_mesh`). Options: pressure-swing (if the azeotrope moves ≥ 5 mol% over the P range), entrainer/extractive, or hybrid with the membrane unit. Set $\Delta T_{\min} \geq 10$ °C on the reboiler/condenser to keep utility cost sane.

4 Absorption, stripping & gas treating

Absorption transfers a gas-phase solute into a liquid solvent; stripping is the reverse. Reach for it instead of distillation when the species is DILUTE in a gas, or when no boil-up is available — you contact, you do not boil.

Hypothesis to run: if the solute is $< \sim 10$ mol% of a gas and a solvent gives a low slope $m = y/x$, absorption is far cheaper than condensing and distilling the whole stream.

Situation	Why absorption/stripping wins
Dilute solute in a large gas stream	no need to condense/boil the bulk inert
Acid-gas removal (CO ₂ , H ₂ S)	selective chemical solvent (amine) near ambient T
Solvent / VOC recovery from a vent	recover below its dew point without refrigeration
Heat-sensitive solute	ambient contacting, no reboiler

The absorption factor and Kremser. The master group is $A = L/(m \cdot V)$ (liquid rate / equilibrium-slope $\times$ gas rate); stripping uses $S = 1/A = mV/L$. Design $A \approx 1.2\text{--}2.0$: at $A < 1$ you can never reach a fixed recovery no matter how tall the tower; at $A > 2$ you waste solvent and the recovery curve flattens. The **Kremser equation** turns A and the target recovery into ideal stages — $\sim 99\%$ recovery at $A = 1.4$ needs roughly **6–10 ideal stages**.

Minimum solvent. The minimum L/V is where the operating line touches the equilibrium line; design at $L = 1.2\text{--}1.5 \times L_{\min}$ — like reflux, less is infeasible and much more just inflates the downstream stripping duty. Always judge the absorber+stripper PAIR, never the absorber alone.

Internals	Pick when	Numbers
Random / structured packing	corrosive, foaming, low ΔP , small diameter ($< \sim 0.6$ m)	HETP 0.3–0.6 m (structured), 0.5–1 m (random)
Trays	large diameter, high liquid load, side-draws, fouling	0.3–0.6 m spacing, 60–80% stage efficiency

Quote packed height as **HTU** $\times$ **NTU** (HTU $\sim 0.3\text{--}1$ m); run the gas at **60–80% of flooding** (the same F-factor/superficial-velocity idea as distillation). **Gas treating:** for CO₂/H₂S use a chemical solvent (MEA/MDEA, 15–50 wt%) — the reaction boosts capacity at low partial pressure, but regeneration costs $\sim 1\text{--}4$ MJ per kg CO₂ and dominates the operating cost, so solvent choice is really a regeneration-energy decision. Cold + high-P favours absorption; hot + low-P favours stripping — the same column run in reverse.

5 Liquid-liquid extraction

Extraction (LLE) separates by *chemical affinity* to a second liquid phase, not by volatility — so it wins exactly where distillation struggles. **Hypothesis to run:** take a feed and check whether the relative volatility α and the solvent distribution coefficient K point the same way; if $\alpha \rightarrow 1$ but K is large, LLE is the cheaper route.

When LLE beats distillation (the four classic triggers):

Trigger	Distillation symptom	Reach for LLE when
Close-boiling	$\alpha < 1.05$ –1.10 $\rightarrow$ 80+ trays, high reflux	a solvent gives $K \geq 5$ –10 with selectivity $\beta > 20$
Azeotrope	pinch, can't cross	solute partitions selectively (aromatics/aliphatics, acetic acid/water)
Heat-sensitive	thermal degradation > 100 –150 °C	ambient-to-40 °C extraction protects the product (antibiotics, vitamins)
Dilute aqueous	boil-up energy $\propto$ whole stream	solute < 5 –10 wt% and $K \cdot (S/F) \gtrsim 5$ (metals, phenol, caprolactam)

The catch: LLE almost always needs a downstream distillation to *recover the solvent*, so judge the **combined** train, not the extractor alone.

Solvent selection — rank candidates on four numbers:

- **Distribution coefficient $K = y/x$** (solute in solvent / in raffinate). Want $K \gtrsim 5$ –10; below ~ 1 the solvent flow balloons.
- **Selectivity $\beta = K_{\text{solute}}/K_{\text{carrier}}$** , the LLE analogue of relative volatility. $\beta > 1$ is mandatory; $\beta > 20$ –50 lets you reach high purity in few stages. $\beta \rightarrow 1$ is the LLE "azeotrope."
- **Recoverability:** a 10–30 °C boiling-point gap (or a density/solubility swing) between solvent and solute so the recovery still is cheap.
- **Density difference $\Delta\rho \geq 50$ –150 kg/m³** and **moderate interfacial tension $\gamma \approx 5$ –30 mN/m** for clean settling — too low γ emulsifies, too high needs intense agitation.

Number of stages. Use the Kremser equation with the extraction factor $E = K \cdot (S/F)$. Design for $E \approx 1.3$ –2.0: at $E < 1$ you can never reach high recovery regardless of stages; at $E > 3$ you are wasting solvent. For 99% solute recovery at $E = 1.5$, Kremser gives roughly **8–12 ideal stages**; raising E to 2 cuts it to ~ 5 –7. Stage efficiencies are low (mixer-settler 75–95%, columns 20–40% per "stage"), so HETS dominates real height.

Solvent-to-feed ratio (S/F). Set by E and K : $S/F = E/K$. With $K = 5$ and $E = 1.5$, $S/F \approx 0.3$; with a weak $K = 1$, the same E forces $S/F \approx 1.5$ — and a big solvent-recovery column. Stay in $S/F \approx 0.2$ –2; outside it the economics flip back toward distillation.

Equipment selection by throughput and settling tendency:

Contactors	Throughput	Stages (typ.)	Use when
Mixer-settler	low–very high	1 per unit, trains of 3–6	slow-settling, high efficiency, many stages, easy scale-up (footprint large)
Sieve-tray / packed column	medium–high	5–10	good $\Delta\rho$ & γ , few–moderate stages, low cost
RDC / pulsed / Karr (agitated)	medium	8–15	low γ , emulsifying systems, tall but compact footprint
Centrifugal (Podbielniak)	low–medium	3–7	tiny $\Delta\rho$, short residence (antibiotics), high cost

Why: agitation buys mass transfer but fights settling — the design tension is always *contact intensity vs. phase disengagement*. In Choupo, sweep S/F and stage count on an absorber/extractor case and watch the Kremser recovery curve flatten past $E \approx 2$ before committing solvent.

6 Membranes & aqueous/advanced separations

Membranes win where the separation is **isothermal and low-grade-energy-only**: aqueous streams, heat-sensitive or non-volatile solutes, and modest concentration ratios. The rule of thumb: **if the relative volatility is favourable and the solvent must be boiled anyway, distillation/evaporation wins; if you only need to concentrate an aqueous solution or strip ions/macromolecules at near-ambient T, a pressure-driven membrane is 5-20× cheaper in energy**. RO desalination runs ~3-4 kWh/m³ vs ~10-15 kWh/m³ (with vapour recompression) for thermal — because you pay only to overcome osmotic pressure (~0.78 bar per 1000 mg/L NaCl), not the latent heat of water.

6.1 Selecting the membrane class

Class	Pore / MWCO	ΔP (bar)	Flux (LMH)	Rejects	Use when
MF	0.1-10 μm	0.5-2	50-1000	suspended solids, bacteria	clarification, cell harvest, RO pretreatment
UF	1-500 kDa	1-5	30-150	macromolecules, colloids, virus	protein concentration/-diafiltration, oily water
NF	200-1000 Da	5-25	15-60	divalents, sugars, dyes; passes monovalents	softening, sugar/whey, colour removal
RO	<200 Da	15-70 (BWRO) 10-25, SWRO 55-70)	12-30	nearly all ions	desalination, final dewatering, ZLD

Driving pressure scales with what you reject: as MWCO falls, osmotic back-pressure and fouling resistance rise, so flux drops and ΔP climbs. NF is the "selective" sweet spot — it passes NaCl (low osmotic penalty) while holding SO₄²⁻/Ca²⁺ (*membrane02_NF_sugar*).

6.2 Recovery vs scaling — the master trade-off

Higher single-pass recovery r shrinks capital and brine volume, but concentrates sparingly-soluble salts at the wall by the **concentration factor CF = 1/(1-r)**. At $r = 0.75$, CF = 4; at $r = 0.90$, CF = 10. Cap recovery so the **scaling saturation index stays below the threshold**: keep **LSI < 0 (or < +1.8 with antiscalant)** for CaCO₃ and **S&DSI / IAP-Ksp ratio < 0.8** for CaSO₄, BaSO₄, SiO₂ (<120 mg/L). This is exactly the recovery-vs-SI sweep in *optim04_membrane_recovery_scaling* and *membrane07_scaling_si* / *membrane08_softened_scaling* (softening lifts the ceiling). Typical r : **SWRO 35-50% per pass, BWRO 75-85%, NF 80-90%** with antiscalant.

6.3 Crossflow, flux and polarisation

Concentration polarisation multiplies wall concentration by **CP = exp(Jw/k)**; keep it < **1.3-1.5**. Push the mass-transfer coefficient k up with crossflow: spiral-wound feed-channel **superficial velocity 0.1-0.4 m/s** (Re 100-1000, spacer-filled), tubular/UF **2-6 m/s**. Design flux conservatively below the **critical/limiting flux**: **SWRO 12-17 LMH, BWRO 18-27, NF 15-40, UF 30-100**. Over-fluxing collapses k and locks in irreversible fouling. *membrane03_concentration_polarization* and *membrane04_spacer_schock_miquel* let you watch CP rise as velocity drops.

6.4 Staging, trains, pretreatment

Spiral-wound elements are 8"×40", ~37 m² each, 6-8 per pressure vessel. To hold velocity as permeate is removed, taper the **array (e.g. 2:1 → 6 then 3 vessels)**; add a **second pass** on permeate when you need <500 mg/L product. Run a tapered train in *membrane05_train*. Budget **per-vessel ΔP 1-3 bar** (replace at 15% rise). Always pretreat: MF/UF or coagulation to **SDI₁₅ < 3**, antiscalant, and cartridge filtration — fouling, not membrane area, sets RO lifetime.

7 Solids processing & particle technology

Solids units are governed by the **particle**, not the molecule: crystal habit, size distribution and mechanical strength dictate every downstream choice. Always carry a target median size (d_{50}) and span as design specs.

7.1 Crystallisation route — pick by the solubility curve

The route follows how solubility varies with the achievable handle. Generate supersaturation ($S = c/c^*$) gently: hold the **metastable zone width**, typically $\Delta c/c \approx 0.01\text{--}0.05^*$ ($S \approx 1.01\text{--}1.05$), so growth dominates nucleation and you get coarse, filterable crystals. Push S past ~ 1.1 and you nucleate a fine, hard-to-filter slurry.

Route	Use when	Driving handle	Typical yield knob	Watch
Cooling	dc^*/dT steep (>1 wt%/10 °C, e.g. KNO_3 , KCl)	ΔT , 0.1–1 °C/min ramp	ΔT over solubility curve	wall encrustation; keep $\Delta T_{\text{wall}} < 5\text{--}10$ °C
Evaporative	flat solubility (NaCl , ~ 0.07 wt%/°C)	solvent boil-off	evaporation rate	scaling, foaming; vacuum to cut T
Antisolvent	high solubility, T -insensitive; heat-sensitive API	miscible non-solvent addition	solvent:antisolvent ratio	local high- $S \rightarrow$ fines; add slowly, well-mixed
Reactive/precip.	sparingly soluble product	reagent feed rate	stoichiometric excess	very fast nucleation $\rightarrow$ nanoparticles

MSMPR sizing: mean residence $\tau = 1\text{--}4$ h; the population balance gives $d_{43} \propto G\tau$, so longer $\tau \Rightarrow$ bigger crystals at fixed growth rate G ($\sim 10^{-8}\text{--}10^{-7}$ m/s). See `crystalliser02_msmpr` (population balance), `crystalliser06_nacl_antisolvent` and `crystalliser07_nacl_antisolvent_cooled` for combined routes.

7.2 Solid–liquid separation — filtration vs centrifugation

Pick by particle size and solids fraction. Vacuum/pressure filtration suits $d_{50} > 10\text{--}20$ μm ; below ~ 5 μm cake resistance explodes (specific cake resistance α rises) and you switch to centrifugation or add filter aid. Sedimenting centrifuges handle $\Delta\rho > 50$ kg/m^3 and feeds too fine or too compressible to filter; decanters take 5–40 wt% solids, disc-stacks < 5 wt%. Rule: filter when cake washing/dryness matters and particles are coarse; centrifuge when fines, fragile crystals, or fast throughput dominate.

7.3 Dryer selection — by feed state and particle size

Dryer	Feed	Product d_{50}	Residence	Notes
Spray	pumpable slurry/solution ($< 40\%$ solids)	10–200 μm	5–30 s	hollow spheres; high T_{gas} but short $\rightarrow$ heat-sensitive OK. <code>sprayDryer01_sugar</code>
Fluid-bed	free-flowing granules	50 μm –5 mm	10–60 min	narrow size range; needs U above U_{mf}
Drum	viscous paste/slurry	flakes	5–30 s	conductive, no carrier gas
Tray/batch	small lots, fragile	any	2–24 h	cheap, labour-heavy

`solidDryer01_sugar` shows a through-circulation solid dryer; `sprayDryer02_residence_sweep` sweeps τ vs moisture.

7.4 Gas–solid separation & fluidisation velocities

Cyclones are the workhorse for $d > 5\text{--}10$ μm ; below that, efficiency falls and you add a bag filter (`bagFilter01_dust`). Design the **inlet velocity to 15–25 m/s** (> 30 m/s re-entrains/erodes; < 10 m/s drops cut size); body-to-inlet $\Delta P \approx 0.5\text{--}2$ kPa, 4–8 inlet velocity heads. See `cyclone01_dust_removal`, `cyclone02_leith_licht`.

Superficial-velocity targets: **fluid-bed operate at $2\text{--}4 \times U_{\text{mf}}$** (above minimum fluidisation, below terminal U_{t} to avoid carryover); **pneumatic conveying dilute-phase 15–25 m/s**; settling-chamber gas < 3 m/s. Always check the actual U_{mf} and U_{t} from particle d and $\Delta\rho$ rather than trusting the ratio blindly — that is the glass-box hypothesis to run.

8 Evaporation & multiple-effect

Evaporation concentrates a solution by boiling off solvent (usually water). The whole game is **steam economy** — kg water evaporated per kg live steam — and the way to raise it is to reuse each effect's vapour as the next effect's heating steam. **Hypothesis to run:** sweep the number of effects on Choupo's triple-effect case and watch economy climb while area and capital grow.

Effects	Steam economy	Pick when
1	0.8–0.95	small duty, cheap steam, scaling/fouling service
2–3	1.7–2.6	medium duty — the common industrial choice
4–6	3.5–5.0	large water removal, expensive steam (sugar, paper, desal)

Each added effect adds roughly **+0.8 kg evaporated per kg steam**. WHY it saturates: the available ΔT (steam T minus the vacuum/condenser T) is **split** across the effects, so each has a smaller driving force and needs more area. Capital $\propto$ effects while the economy gain shrinks — the optimum is usually **2–5 effects**.

Feed arrangement	Direction	Use when
Forward	feed with steam	feed near boiling; simplest, no inter-effect pumps
Backward	feed counter to steam	cold or viscous feed (the hottest effect handles the thickest, lowest-viscosity liquor)
Mixed	hybrid	balance viscosity and economy

Boiling-point elevation (BPE): dissolved solids raise the boiling point above pure water — a few °C when dilute, **15–25 °C** for strong NaOH or brine — eating into the ΔT and capping the last effect (Choupo's NaOH and sugar evaporator cases account for it). **Type selection:** falling-film for low- ΔT , heat-sensitive, low-viscosity service (high U); **forced-circulation** (pumped 2–3 m/s) for viscous, scaling, or crystallising liquors; agitated/wiped-film for very viscous. **Vapour recompression** (electric MVR, or TVR with a steam ejector) recycles the vapour's latent heat into the same effect — one MVR effect can match a 4–6-effect train when electricity is cheap relative to steam.

9 Gas-liquid separators & process drums

Every flash, every reflux accumulator, every compressor suction line needs a vessel that lets vapour and liquid part ways. The physics is a settling race: a liquid droplet falls under gravity while the rising vapour drags it up. Size the vessel so the vapour goes slow enough that droplets above a cut size ($\sim 150\text{ }\mu\text{m}$ bare, $\sim 10\text{ }\mu\text{m}$ with a mesh pad) rain out instead of carrying over.

The Souders-Brown anchor. The maximum allowable vapour velocity is $v_{\text{max}} = K \cdot \sqrt{(\rho_L - \rho_V) / \rho_V}$. The square-root group is just the terminal-settling balance (gravity vs. drag); K (m/s) bundles droplet size, drag coefficient and a safety margin. The cross-section follows from $A = Q_V / v_{\text{max}}$, so K sets the diameter directly — get it wrong and you either spray liquid downstream or buy steel you didn't need.

K-factor (m/s)	Service / condition	Why
0.03–0.06	No demister, bare gravity drum	Only $\sim 150\text{ }\mu\text{m}+$ drops settle; derate hard
0.06–0.07	Vacuum / low P (< 1 bar)	Low ρ_V makes the group large; cap velocity
0.08–0.11	Mesh demister pad, clean service	Pad coalesces $5\text{--}10\text{ }\mu\text{m}$ drops; run faster
$\times 0.7\text{--}0.8$	Foaming / fouling / glycol-amine	Foam re-entrains; cut K by 20–30 %
$\times 0.9$ per	0.5–1.0 m of disengagement below pad	Short disengaging height steals settling time

Use the **vapour density at operating T,P**, not standard — a compressor suction drum at 5 bar has $\sim 5\times$ the ρ_V of atmospheric, shrinking v_{max} by $\sim 2.2\times$ and demanding more area than newcomers expect.

Vertical vs. horizontal. Choose by vapour-to-liquid loading and footprint.

Pick	When	Why
Vertical	High vapour / low liquid; compressor suction & K.O. drums ; tight plot	Settling distance = height is cheap to add; small liquid hold
Horizontal	High liquid load, large surge/holdup, reflux accumulators , slug catchers, 2 liquid phases	Big interfacial area per volume; long L for L-L decant; lower vessel

For horizontal drums the disengagement area is $L \cdot (D \text{ minus liquid level})$, so set the vapour space to $\sim 40\text{--}50\%$ of the cross-section and check v_{max} against that reduced area.

Geometry and holdup. Aim for $L/D \approx 3\text{--}5$ (below 2 the heads dominate cost; above 6 it gets whippy and hard to support). Liquid **residence time 5–15 min** of working volume sets the length/level: ~ 5 min for a clean reflux drum feeding a good level controller, 10–15 min where downstream is a fired heater or compressor you must never starve or slug. Allow **vapour disengagement height $\approx 0.3 \cdot D$ or 0.45 m (whichever larger)** between the high liquid level and the inlet/demister, plus $\sim 0.15\text{ m}$ mesh pad. Inlet nozzle momentum $\rho \cdot v^2$ should stay $< 2250\text{ kg/m} \cdot \text{s}^2$ ($< \sim 1000$ if no inlet device) or the jet shatters liquid faster than the pad can catch it.

Glass-box hypotheses to run. In Choupo a flash drum is just an `IsothermalFlash` at (T,P); the vapour and liquid rates it returns are your Q_V and holdup feed. Try `tutorials/steady/flash/flash01_benzene_toluene`: take the V and L rates, pick $K = 0.08$ (pad) vs. 0.04 (bare), and watch the required diameter roughly double when you drop the demister and halve K . Then re-flash at 5 bar and confirm the higher ρ_V forces a smaller v_{max} and a fatter drum — the Souders-Brown sqrt made visible.

10 Heat exchangers — type selection & sizing

Pick the *type* first — it sets cost, U , and ΔP before you ever size area. The driving question is the **overall coefficient** U : a gas film ($h \approx 30\text{--}300\text{ W/m}^2\text{K}$) starves U , so the cheap-per- m^2 geometry wins; a clean liquid–liquid duty ($h \approx 1000\text{--}5000$) lets you chase compactness and low $\Delta T_{\min}$.

10.1 Type-selection table

Type	Best service	U ($\text{W/m}^2\text{K}$)	ΔP budget	T/P limit	Why
Shell-and-tube (TEMA)	Any phase, high P/T, fouling	300–1200 (liq-liq), 30–300 (gas)	0.3–0.7 bar	<40 bar, <600 °C	Rugged, mechanically cleanable, ASME default
Plate-and-frame (gasketed)	Clean liquid–liquid, close approach	2000–7000	0.3–1 bar	<25 bar, <180 °C	Huge area/volume, $\Delta T_{\min} \approx 1\text{--}5$ °C; gaskets limit T
Air-cooled (fin-fan)	Hot streams, no cooling water	15–60 (bare-tube basis)	gas side ~ 0.005 bar	tube-side any	Saves water; approach to ambient $\geq 10\text{--}15$ °C
Double-pipe	Small duty (<20 m^2), high P, dirty	300–1000	flexible	very high P	Cheap, true counter-current ($F=1$)
Spiral	Slurries, fouling, viscous	700–2500	self-scouring	<20 bar	Single channel resists plugging
Printed-circuit (PCHE)	Very high P, compact (LNG, H_2)	1000–5000	tight	<600 bar	Diffusion-bonded; near-zero fouling tolerance

10.2 Gas-phase service

Shell-and-tube *still* makes sense for gas — but put the **gas on the shell side** with the gas film controlling, accept $U \approx 30\text{--}300$, and expect large area. For sensible gas cooling against ambient, **air-cooled** usually beats shell-and-tube + cooling-water tower on TAC once duty > ~ 0.5 MW. For gas–gas (e.g. feed/effluent), a plate-fin or PCHE recovers heat that a shell-and-tube cannot economically reach.

10.3 Allocation & velocities

Put on the **tube side**: the dirtier, higher-pressure, more-corrosive, or lower-flow stream (easier to clean, cheaper to alloy small tubes).

- Tube-side liquid: **1–2.5 m/s** (>1 m/s suppresses fouling; <3 to limit erosion; <1.5 for SS to avoid erosion-corrosion).
- Shell-side liquid (cross-flow): **0.3–1 m/s**.
- Gas, tube or shell: **10–30 m/s** (process gas), up to 30 m/s for low-density vapour; chase the gas film, watch ΔP .
- Condensing vapour: keep below $\sim 50\%$ of erosion velocity; flooding governs.

10.4 Phase-change service

- **Condensers**: prefer shell-side condensation (vapour on shell, coolant in tubes); $U \approx 800\text{--}1500$ for organics, 1500–4000 with steam. Total condenser $\rightarrow \Delta T$ driven; allow vapour velocity de-rating.
- **Reboilers**: kettle ($U \approx 800\text{--}1200$) for wide turndown and dirty bottoms; thermosiphon ($U \approx 1000\text{--}1500$) cheaper but needs $\Delta T >$ a few °C and stable circulation. Cap heat flux below the **critical $\sim 30\text{--}40\text{ kW/m}^2$** to dodge film boiling. See [Choupo tutorials/steady/distillation](#) (condenser/reboiler duties) and [tutorials/steady/heat](#).

10.5 $\Delta T_{\min}$, LMTD correction, fouling

- **$\Delta T_{\min}$:** 10–20 °C for hot utility/gas, 5–10 °C liquid–liquid, 3–5 °C refrigeration (energy vs area trade-off — verify via Choupo pinch composite curves).
- **LMTD correction F :** keep **$F > 0.8$** ; $F < 0.75$ means add shell passes (2-4, 4-8) or go pure counter-current.
- **Fouling:** add $R_f \approx 0.0002$ (clean liquid) to 0.0005–0.001 m²K/W (cooling water, fouling organics); this can dominate $1/U$ and is why a glass-box run should report the clean-vs-fouled area split.

11Heat integration & pinch analysis

Heat-exchanger SELECTION sizes one match; pinch analysis decides which streams should match AT ALL and sets the energy target before any network is drawn. It is the highest-leverage utility-saving tool — typically **20–30% off the energy bill** versus a naive design. **Hypothesis to run:** build the composite curves in Choupo's Pinch view and read the targets before placing a single exchanger.

Composite curves and the pinch. Combine all hot streams into one composite (T vs $\Sigma C_p \cdot \Delta T$) and all cold streams into another; slide them together until the closest vertical approach equals ΔT_{min} . That point is the **PINCH**; the overshoots are the minimum hot and cold utilities (the **MER targets**), fixed before design.

Service	ΔT_{min}	Why
Liquid–liquid	10 °C	default; balanced area vs energy
Process with a gas stream	20–40 °C	low $U \rightarrow$ area explodes at small ΔT
Refrigerated / cryogenic	3–5 °C	cold is expensive; pay area to save it
Near a furnace / high-grade heat	20–60 °C	high-grade heat is cheap relative to area

Supertargeting plots total cost vs ΔT_{min} and picks the flat minimum (often **8–15 °C** for fluids).

The three golden rules — break any and you spend MORE than the target:

1. **No heat across the pinch** — a cross-pinch exchanger adds the same amount to BOTH utility bills.
2. **No hot utility below the pinch** — below it there is a heat surplus; adding heat just needs more cooling.
3. **No cold utility above the pinch** — above it there is a deficit; cooling there just needs more heating.

The grand composite curve (GCC) plots surplus/deficit vs shifted T : read it to place utilities at the RIGHT level — pick the cheapest steam/refrigerant the pockets allow, spot an intermediate utility (cheaper MP steam, or raising steam from recovered heat), and see where a fired heater or molten salt is GENUINELY needed versus where recovered heat suffices.

12 Fired heaters & furnaces

Reach for a fired heater when the duty is large (>3 MW) or the process temperature is above what your hottest utility can deliver. HP steam tops out near 250 °C and 40 bar; HiTec molten salt reaches ~535 °C; Dowtherm A ~400 °C. Above that — crude-unit charge heaters (350-400 °C), reformer feed, vacuum-tower heaters, reboilers on high-boiling bottoms — you must burn fuel directly, because the WHY is simple: a flame radiates at 1500-1900 °C, so it can drive heat across a tube wall that no recirculated fluid can. Below ~535 °C and a few MW, molten salt or fired-salt loops usually win on safety and turndown; don't fire directly if a utility reaches the target.

Two heat-transfer zones, by mechanism. The RADIANT section (firebox) absorbs 60-75% of the duty by flame radiation; tube-wall view factor and flame emissivity dominate, so the design knob is *flux*, not area. The CONVECTION section above it scavenges 25-40% of the duty from 700-900 °C flue gas by convection over finned/studded tubes, dropping stack gas to 150-250 °C. The split exists because radiant flux is capped: push too hard and you coke the film.

Heater type	Duty range (MW)	Typical use	Radiant flux (kW/m ²)	Why pick it
Vertical-cylindrical	0.5-20	Small/medium, plot-limited	25-35	Cheapest, smallest footprint, uniform firing
Box (cabin), horizontal tubes	10-100	Crude/vacuum charge, reboil	30-40	Big duty, even flux, many burners
Box, two-fired cell	50-300+	Reformers, ethylene cracking	35-50 (cracking, short residence)	Highest duty, double-fired uniformity
Cylindrical helical-coil	0.5-10	Hot-oil / salt heaters	20-30	Compact, low-fouling clean service

Radiant flux — the master rule. Design average radiant flux 30-40 kW/m² (some references 25-45). Go higher and the *peak* local flux — typically 1.5-2x the average — drives the tube-wall film over its coking threshold (~430 °C for crude oil, lower for residue), laying down coke that insulates, raises tube metal temperature, and shortens runlength. Average flux is therefore a runlength decision, not just an area decision: halving flux can double time-between-decoke.

Efficiency, excess air, draft. Thermal efficiency 75-82% bare; 85-92% with an air preheater recovering stack heat. Each 20 °C of stack-gas drop is worth ~1% efficiency, and 1% excess-air reduction another ~0.5%. Run 10-15% excess air on gas, 15-25% on oil — enough to complete combustion, not so much you heat nitrogen up the stack. Draft: natural (buoyant stack, simplest, cheap) for small heaters; forced or balanced draft once you add an air preheater (the preheater's pressure drop demands a fan). Keep a slight negative draft (-2 to -5 mm H₂O) at the arch so flue gas leaks in, not flame out.

Tube metallurgy sets the limit. Specify tubes by metal temperature = process T + flux-driven ΔT (50-100 °C) + fouling margin: carbon steel to ~400 °C, 1.25Cr-0.5Mo to ~550 °C, 9Cr-1Mo to ~650 °C, HK/HP cast austenitics to 1000 °C+ for reformers. NO_x scales with flame temperature — low-NO_x burners (staged air/fuel, flue-gas recirculation) cut it to 30-50 ppm but lengthen the flame, so check flame-to-tube clearance. Hypothesis to run in Choupo: model the heater as a W-equivalent duty against a fired-fuel stream, then sweep excess air and stack temperature to watch efficiency and fuel firing rate move — confirm each 20 °C stack drop buys ~1%.

Why the tubes do not melt at a 1900 °C flame. Students reasonably ask how steel survives a flame far above its melting point. The answer: **the tube never sees the flame temperature.** The process fluid races through the inside and carries the heat away continuously, so the tube wall sits at roughly **process-T plus a small film rise** — typically **400–600 °C metal**, not 1900 °C. The radiant flux (30–40 kW/m²) is balanced by the inside film: the film rise = flux / h_{inside} is only **30–120 °C** for a fast liquid ($h \approx 1000\text{--}3000$ W/m²K). The flame's heat is real, but it is removed as fast as it arrives.

Two lessons: (1) the metallurgy is specified for the **tube-metal temperature**, not the flame (carbon steel ~400 °C, 9Cr-1Mo ~650 °C, cast HK/HP austenitics to ~1000 °C); the firebox walls are protected by **refractory** lining, not bare metal. (2) If the process flow STOPS or drops (pump trip, coking), the heat is no longer removed and the tube temperature runs away in **seconds** — the tube ruptures. That is why a fired heater carries a **low-flow trip** and the flux is capped: a hot tube with no flow is the classic furnace failure.



Figure 1: Why furnace tubes survive a 1900°C flame. (a) The radiant tubes hang BARE along the firebox walls, exposed to the flame; the refractory lines the WALLS (protecting the steel shell and re-radiating heat), not the tubes. (b) Almost the whole temperature drop falls across the outer gas film, so the tube metal sits near 520°C – far below its melting point – because the process fluid inside removes the heat as fast as it arrives. Lose the flow and the metal races toward the flame temperature.

13 Compressors, expanders & pressure change

Gas pressure change is governed by *volumetric flow at suction* and *required pressure ratio* $r = P_{\text{out}}/P_{\text{in}}$ — not mass flow. Because gas density falls as you compress, machine choice tracks **actual inlet m^3/h** and r , then degrades to the next type as flow rises or r falls. Treat every heuristic below as a hypothesis: run the case in Choupo, read W_{shaft} and discharge T, and check you stayed inside the box.

13.1 TYPE SELECTION

Type	Inlet flow (actual m^3/h)	r per body	Typical η_{poly}	Why / physics
Fan / blower	10^3 – 10^6	<1.1 (fan), <2 (blower)	0.55–0.75	Low Δp ; cheap; r limited before it becomes a "compressor"
Centrifugal	1,700–340,000	1.5–3 per wheel, ~10 multi-stage	0.75–0.85	Dynamic; smooth, oil-free, high flow; surge limits turndown
Axial	>65,000	1.2–1.5 per stage	0.80–0.88	Highest flow, best η ; narrow operating window
Reciprocating	<16,000	up to 5–10 per stage	0.75–0.90	Positive-displacement; best for high r , low flow, high P (>200 bar)
Rotary screw	<5,000	3–4 (up to ~8 oil-flooded)	0.70–0.78	Robust, dirty/wet gas, modest r ; oil carryover

13.2 Stages, ratio per stage & intercooling

Split r across N stages with **equal ratio per stage** $r_{\text{stage}} = r^{(1/N)}$, kept to ~2–4 (cap ~4 for recips, ~3 for centrifugals). Equal ratios minimise total shaft work and equalise per-stage discharge T. The driver of staging is the **discharge-temperature limit**: keep $T_{\text{out}} \lesssim 150$ – 175 °C (lube-oil coking, gasket and rod-packing life; ~200 °C absolute ceiling). Adiabatic $T_{\text{out}} = T_{\text{in}} \cdot r_{\text{stage}}^{((\gamma-1)/(\gamma\eta))}$, so a single jump from 1→16 bar on a $\gamma \approx 1.3$ gas blows past 250 °C — unacceptable. **Intercool back to ~40 °C** between stages: this cuts the next stage's volumetric work (you compress colder, denser gas) and typically saves 10–20 % total power for $r > \sim 5$. Pair an intercooler with a knock-out drum — cooling condenses water/heavy ends that would erode the next stage. In Choupo, model each stage as a **compressor** ($W_{\text{shaft}} + \eta$) feeding a cooler + flash; cooling-water or chilled-water utilities by T-level supply the intercooler duty.

13.3 Efficiency, surge & turndown

Quote **polytropic** efficiency (0.75–0.85 centrifugal, up to 0.88 axial) for design — it is path-independent and additive across stages, whereas isentropic η drifts with r . Multiply by ~0.95 mechanical to get brake power. Dynamic machines **surge** at ~60–70 % of design flow (flow reversal, violent oscillation), so usable **turndown is only to ~70 %** without recycle/anti-surge; recips turn down to ~25 % via clearance pockets or speed. Size with ~10 % flow margin above the surge line.

13.4 When an expander pays back

If a hot, high-P stream must be let down (≥ 3 –4 bar drop, $\geq$ a few hundred kW recoverable), replace the throttle valve with a **turbo-expander** (model as **turbine**, $W_{\text{shaft}} + \eta \approx 0.80$ – 0.85). Recovered shaft work directly drives a generator or a companion compressor (common in refrigeration, FCC flue-gas, cryogenic N_2). The trade-off: expansion chills the stream (Joule–Thomson + work extraction), so verify you don't cross a dew/freeze point — a glass-box run shows the outlet T before you commit hardware.

14 Drivers: electric motors, steam & gas turbines

A pump or compressor is only as available as the machine that turns it. The *driver* is a separate selection from the hydraulic/thermodynamic duty: pick it by shaft power, by what HP steam or fuel gas the plant already has, and by what must keep running when the grid trips. In Choupo the rotating unit carries $W_{\text{shaft}} + \epsilon$; the driver question is *where that shaft power comes from* and at what reliability.

Rule — electric motor is the default below ~0.5–1 MW. Standard TEFC induction motors are cheap, compact, ~93–97% efficient, and need no auxiliaries. WHY the ceiling: above ~1 MW you move to medium-voltage (4.16 kV+) switchgear and the incremental $/kW \text{ of motor} + \text{variable} - \text{frequency drive (VFD) climbs, while steam turbine that} * \text{also} * \text{lets HP steam down to cause full level start to win on combined steam} - \text{balance economics. Hypothesis to run} : \text{size the motor at } 110 - 1250.06 - 0.10 / \text{kWh}.$

Rule — steam turbine when HP steam exists, for very large duty, or as a trip-proof spare. A back-pressure turbine taking HP→MP/LP steam "pays" for shaft work by letting steam down a pressure level you needed anyway — marginal cost is only the lost power-generation credit, so isentropic efficiency of 65–80% (small) to ~85% (large) is acceptable. A condensing turbine (exhaust to ~0.1 bar) maximises work but dumps ~2 MJ/kWh of latent heat to cooling water. Critically, a steam turbine keeps spinning on a grid blackout, so the boiler-feedwater pump and reactor-quench pump are classically *steam-turbine-driven spares* on a motor-driven main (or vice-versa).

Rule — gas turbine only for very large power + heat (combined heat-and-power). Below ~5 MW the heat-rate and capital rarely justify it; above ~10–15 MW a gas turbine with a heat-recovery steam generator delivers shaft power *and* the HP steam header — the cogeneration case Choupo models with steam utilities by T-level.

Driver	Sweet-spot power	Efficiency	Pick it WHEN	Watch-out
Electric motor (fixed-speed)	< 0.5 MW	93–97%	small/medium duty, throttled or constant flow	dumps control margin across a valve
Electric motor + VFD	0.1–5 MW	90–95% (incl. drive)	flow varies > ±20%	drive cost, harmonics, ~3–5% drive loss
Back-pressure steam turbine	0.5–15 MW	65–85% isentropic	HP steam available, want MP/LP let-down	power tied to steam demand
Condensing steam turbine	1–30 MW	70–85% isentropic	large duty, run-on-trip spare	condenser + cooling-water latent load
Gas turbine (CHP)	> 10 MW	30–40% (≈80% w/ HRSG)	huge power + process steam	fuel-gas spec, NOx, hot-day derate

Rule — variable speed beats throttling when flow turndown exceeds ~±20%. Power scales ~speed³ on a centrifugal machine, so trimming flow by slowing the driver saves far more than burning the surplus head across a control valve. WHY not always: a VFD adds ~3–5% standing loss and capital, so for near-constant flow a fixed-speed motor + control valve is cheaper and simpler.

Rule — spare to n+1 on critical service. Install one extra parallel machine (e.g. 3×50% pumps, run 2) so a single failure or maintenance outage does not stop the plant; mix drivers (one motor, one turbine) so a *common-mode* loss — grid or steam — still leaves one runner. Verify by toggling a unit offline and re-running the flowsheet: does the remaining capacity still meet duty?

15 Fluid transport, line sizing & hydraulics

Moving fluid costs capital (pipe + driver) and power (friction). Every rule below trades a *bigger pipe / lower velocity* (more steel, less ΔP , less erosion) against a *smaller pipe / higher velocity* (cheap, but pumping power and noise climb as $\sim v^2 \cdot s^8$). Pick the driver by **pressure-ratio first, then flow**.

15.1 Driver selection by pressure ratio

Service	Δp / ratio	Pick	Why (glass-box)
Incompressible liquid	any head	Pump (centrifugal <~3500 m, recip/PD for low-flow high-head, viscous, or metering)	$\rho \approx \text{const} \rightarrow$ head, not ratio, is the variable
Gas, $\Delta P < 0.03$ bar	ratio ≈ 1.0 – 1.03	Fan	ρ change negligible, treat as incompressible
Gas, 0.03 – 3 bar	ratio ~ 1.03 – 1.3	Blower (one stage)	mild heating, no intercool
Gas, ratio $> \sim 1.3$	up to ~ 3 – 4 /stage	Compressor (centrifugal; recip for very high ratio/low flow)	T_{out} rises with $\text{ratio}^{((\gamma-1)/\gamma)} \rightarrow$ stage + intercool if total ratio $> \sim 4$

Multi-stage rule: keep per-stage ratio $\approx (\text{total ratio})^{(1/N)}$ so each stage does equal work and discharge T stays below ~ 150 – 200 °C (lubricant/material limit). See `tutorials/steady/rotating/compressor01_air` and the multi-stage logic behind `compressorTurbine01_turbo_booster`; pumps in `pump01_water` / `pump02_pressure_spec`.

15.2 Recommended superficial velocities

Line / service	v (m/s)	Note
Pump suction (liquid)	0.5–1.5	low to protect NPSH; $\Delta P_{\text{suction}}$ is precious
Pump discharge / process liquid	1.5–3 (up to ~ 4.5)	economic optimum; >4 erosion-corrosion risk
Liquid, viscous/slurry	0.5–1.5	keep below settling; above abrasion limit
Gas / vapour, atmospheric–low P	15–30	density-limited
Gas, high pressure	10–20	ρ high $\rightarrow$ same ρv^2 at lower v
Steam (sat./superheat)	20–40 (to 60)	erosion + wet-steam wire-drawing
Two-phase / flashing	10–25	avoid slug/annular; check $\rho v^2 < \sim 6000 \text{ kg/m} \cdot \text{s}^2$

The economic optimum tracks **constant** ρv^2 (kinetic head), which is why dense fluids run slower. A quick screen: $v_{\text{econ}} \approx 2 \text{ m/s liquid, } 20 \text{ m/s gas}$, then verify ΔP in `tutorials/steady/hydraulics/pipe01_water_line`.

15.3 Economic pipe diameter & ΔP budget

- Rough optimum: $D_{\text{econ}} \approx 0.26 \cdot \dot{m}^{0.5} \cdot \rho^{-0.37}$ (m, SI) for turbulent steel — the balance of pipe capital vs pumping power.
- ΔP budget:** liquid lines **0.2–0.5 bar per 100 m** ($\approx 0.5 \text{ kPa/m}$); gas lines **0.1–0.2 kPa/m**; long pipelines push lower (0.05). Total pump head = static lift + friction + equipment + **control-valve allowance**.
- Control-valve ΔP :** allocate the *larger* of **20–33 % of dynamic (friction) ΔP** or **~ 0.7 bar (10 psi)**, so the valve keeps authority at high flow. Too little $\rightarrow$ loop goes open-loop; too much $\rightarrow$ wasted head.

15.4 NPSH (the pump killer)

Require $NPSH_{\text{available}} - NPSH_{\text{required}} \geq 0.5$ – 1 m margin. $NPSH_a = (P_{\text{suction}} - P_{\text{vap}})/\rho g$ + static head – suction friction. Levers: raise the suction vessel, fatten/shorten suction line (low v , see table), subcool, or cut speed. Hot/near-boiling liquids (reboiler bottoms, deaerator) are the classic cavitation traps — design suction generously.

Always run it: these are *hypotheses* — size the line, then let Choupo print the actual ΔP , T_{out} and head so the chosen velocity and driver prove themselves.

16 Control valves, pressure letdown & choking

A control valve is a deliberate, adjustable restriction — you spend pressure to gain control. Two questions fix the hardware: which valve body and trim, and whether the drop is so large it **CHOKES**.

Valve	Rangeability	Use when
Globe (single / multi-stage trim)	20–50:1	throttling — the default; high ΔP ; anti-cavitation / low-noise trim available
Ball / segmented	100–300:1	high rangeability, slurries, on-off-to-throttle
Butterfly	20–50:1	large lines, low ΔP , cheap

Pick the **inherent characteristic** to linearise the loop: **equal-percentage** when most of the system ΔP is elsewhere (the common default — valve gain rises with flow to offset its falling ΔP); **linear** when valve ΔP dominates; **quick-opening** for on-off. Liquid sizing (non-flashing): $Q = C_v \cdot \sqrt{\Delta P / SG}$; size for ~80% open at max flow and ~20% at min, so the plug never sits on the seat (no control) or wide open (no margin).

Choking — the answer to "is a big letdown done in one stage?"




Figure 2: Gas flow through a valve vs the pressure ratio. Below the critical ratio (≈ 0.53 for air) the flow chokes – sonic at the vena contracta – and mass flow stops rising as P_2 falls; the surplus pressure drop becomes noise and erosion.

For a GAS, flow does NOT keep rising as you drop the downstream pressure. Once the pressure ratio falls below the **critical ratio** — $r_c = [2/(\gamma+1)]$ raised to $[\gamma/(\gamma-1)]$, ≈ 0.5 for most gases (0.53 for air) — the velocity at the vena contracta reaches the speed of sound and the valve **CHOKES**: mass flow saturates and the surplus ΔP is dumped as **noise (often > 110 dB), vibration and erosion**. Discharging 50 bar $\rightarrow$ 1 atm is a ratio of 0.02, i.e. deeply choked.

So a single standard valve CAN pass the flow, but across a deeply choked drop it screams, erodes its trim and shakes the pipe — which is why a large letdown is usually **staged**:

Letdown situation	Do this
Small gas ΔP , sub-critical	one globe valve
Large gas ΔP , choked	multi-stage / labyrinth trim, or valves in series, + a downstream silencer (each stage takes a sub-critical drop)
Liquid, $P_2 > P_{vap}$	standard trim
Liquid, $P_2 < P_{vap}$ (flashing)	the liquid flashes/cavitates and imploding bubbles pit the trim — use anti-cavitation (multi-step) trim, harden it, or flash on purpose into a drum

Rule of thumb: stage it when the predicted noise exceeds ~85–90 dBA or the trim-exit velocity is high. The glass-box check: drop the pressure across a **valve** in Choupo and read the outlet T and phase — a flashing liquid shows up as a vapour fraction at the outlet before you ever buy the trim.

17 Refrigeration & vacuum systems

Cold and vacuum are *bought power*. Both reject heat or pull gas against the ambient, so the design instinct is the same: **climb only as far down the ladder as the duty forces you, because every rung multiplies the shaft work.**

17.1 Refrigeration — pick the warmest level that works

Cooling water (CW) leaves a tower at ~30 C and can chill a process stream to ~35-40 C (10 C approach). The moment you need a process temperature **below ~40 C you have left the free zone** and pay for a compressor. The reason is the reverse-Carnot ceiling: ideal COP = $T_{\text{cold}}/(T_{\text{hot}}-T_{\text{cold}})$. Reject to 40 C cooling water and the COP collapses as the cold level drops — so shaft work per unit of cooling *rises* the colder you go.

Cold level needed	Medium	Refrig. T (C)	Ideal COP (reject ~40 C)	Real shaft kW per kW cooling
> 40 C	Cooling water	n/a	—	~0 (pump only)
5-15 C	Chilled water	~5	~6	0.15-0.25 (4-7 kW/kW)
–40 to 0 C	Ammonia / propane	–35	~2.5-3	0.30-0.45 (2-3 kW/kW)
–70 to –40 C	Propylene	–45	~2	0.45-0.55
–100 to –70 C	Ethylene (cascade)	–90	~1.3	0.6-0.9 (~1.3 kW/kW)
< –100 C	Ethylene/methane cascade	–150	< 1	> 1

Rule of thumb: budget ~1 kW shaft per 3-5 kW of cooling at the –40 to 0 C levels, climbing past ~1 kW/kW below –100 C. A single-stage cycle is fine for a ~40-50 C lift; beyond that, **cascade** (ethylene rejecting into propylene rejecting into CW) because one compressor over a 5:1+ pressure ratio runs hot and inefficient. In Choupo the refrigeration compressor is just a **compressor** ($W_{\text{shaft}} + \eta$) feeding the cold side of an exchanger; the warm reject side is **coolingWater**, so the power penalty appears explicitly in the work ledger — model it, don't assume it.

When to refrigerate (not before): a low-temperature condenser to recover a light product (C3/C2 overhead), a cold separator to hit a vapour-loss spec, or a reactor that must run cold for selectivity. Always ask first whether **higher pressure** buys the same condensation at a CW level — compressing the overhead is often cheaper than chilling it.

17.2 Vacuum — match the pump to the absolute pressure

Vacuum distillation earns its keep when the bottoms would **thermally crack or polymerise** at the atmospheric boiling point (heat-sensitive: glycols, fatty acids, monomers, vacuum gas oil). Dropping to **0.05-0.1 bar (40-75 mmHg)** can cut boiling points by 80-150 C, keeping the reboiler below the degradation limit and off the high-T utilities (Dowtherm A / fired heaters). Below ~0.1 bar the **column diameter balloons** (vapour density falls, so allowable mass velocity drops); use low-pressure-drop structured packing, not trays.

Absolute pressure	Technology	Notes / WHY
1-0.1 bar	Liquid-ring pump	Rugged, tolerates condensables; service water sets the floor (~30-50 mmHg)
0.1-0.01 bar	Single-stage steam ejector	No moving parts, uses MP/HP steam; cheap but steam-hungry
< 0.01 bar (to ~1 mmHg)	2-3 stage ejectors w/ intercondensers	Stages compound the ratio; condensers cut steam load
Clean, deep, dry	Dry screw pump	Electric, no steam/water effluent; higher capex, watch for fouling

Size ejector **motive steam by the air/leakage load**, not the process vapour — a few kg/h of inleak through flanges sets the bill, so a tight column beats a big ejector.

18 Storage & intermediate tanks

Storage is inventory bought with steel: too little and a late truck or a tripped upstream unit starves the plant; too much and you pay for tankage, land, and a larger fire/spill exposure. Size by **days of stock** tied to the *logistics* of each stream, then pick a roof that matches the volatility. Run the hypothesis in Choupo by feeding the tank's design flow into a flash drum at storage T,P and watching the vapour fraction — if $vf > 0$ you have a breathing/venting problem a cone roof cannot hold.

Sizing by days of stock (the WHY = decoupling time-scale):

- **Raw materials: ~7-30 days.** The driver is delivery reliability and shipment size — rail/marine (large parcels, long lead) push to 20-30 days; reliable pipeline or daily truck supply drops to 3-7. Hold enough that a missed delivery does not stop the reactor.
- **Products: ~7-30 days,** set by shipping cadence and demand swing. A monthly tanker means ~30 days plus working heel; spot-truck sales need only a few days.
- **Intermediate / surge: hours to ~2 days.** Purpose is to *decouple* two units so a transient in one does not propagate. Rule of thumb: **5-15 min** of holdup for pump-suction/control surge, **0.5-2 days** between batch and continuous sections (e.g. a reactor feeding a continuous still). More surge = looser control coupling but slower composition response and more off-spec inventory if upstream drifts.

Tank **working volume** = days $\times$ daily throughput, then divide by the **fill fraction**: never design to the brim. **Fill ~85-90%**, leaving **10-15% ullage** for thermal expansion (a 30 °C swing on a hydrocarbon adds ~3-4 vol%), foam, and high-level-alarm margin. Overfill is a top cause of major releases (Buncefield), so the ullage is a safety budget, not slack.

Type selection (volatility & pressure drive the roof):

Tank type	Best for	Pressure / vapour-pressure range	Typical size	Why
Cone (fixed) roof	Non-volatile liquids: water, heavy oils, glycols, most products with RVP < ~10 kPa	Atmospheric, near-zero gauge	up to ~10,000 m ³	Cheapest; vapour space breathes in/out — fine when there's little vapour to lose
External floating roof	Volatile liquids: gasoline, naphtha, light solvents (RVP ~10-80 kPa)	Atmospheric	~1,000-100,000 m ³	Roof rides on the liquid → no vapour space , cuts breathing/standing losses by ~95% vs cone roof; emissions + fire-risk reduction
Internal floating roof	Volatiles needing weather/rain protection	Atmospheric	medium	Floating deck under a fixed roof — losses control plus debris exclusion
Sphere (Hortonsphere)	LPG, propane, butane, light ends	~2-18 barg (vapour pressure > atm)	~1,000-10,000 m ³	Sphere carries pressure with minimum steel/area; even stress distribution
Bullet (horizontal cylinder)	Smaller pressurised stock: LPG, NH ₃ , light ends	~5-25 barg	< ~500 m ³	Cheaper than a sphere below ~200-500 m ³ ; easy to transport and mound

The crossover rule: if the stored fluid's true vapour pressure exceeds atmospheric at the *hottest* expected ambient (summer-skin ~45 °C), it **cannot** sit in an atmospheric tank — it must go pressurised (sphere/bullet). Between non-volatile and pressurised lives the floating-roof band, chosen to cut working+breathing losses you would otherwise vent or send to a flare/recovery unit. Always state the range and let throughput, RVP, and emissions cost decide — there is no single right tank.

19 Process control & instrumentation heuristics

A flowsheet that balances on paper still must hold steady against disturbances. The Pareto question is not tuning — it is **what to pair with what**.

What Choupo can and cannot check for you, stated before you rely on it. An earlier version of this paragraph invited you to “step a disturbance in `ctrl01` and watch whether your loop rejects it or fights another”. That invitation was false three ways, and it is corrected here rather than quietly deleted, because a manual that has promised something owes the reader the correction: `ctrl01` carries no disturbance at all (it is a setpoint-tracking case — the disturbance cases are `ctrl02`, `ctrl05`, `ctrl09`); *no* case in the corpus carries more than one PID loop, so “fights another” cannot be watched anywhere; and this section is about *column* pairing, while the dynamic tier has no column in it. A pairing for a column cannot be tested dynamically in Choupo today.

What you *can* do, and it answers the sharpest question in this section — *which tray carries the temperature measurement* — is read it off the rigorous column you already ran. The steady profile gives the slope criterion directly, and re-solving with a manipulated variable nudged gives the sensitivity criterion. The EduTools *Column control* tool does exactly that (§??) and reports something a single answer would hide: on a 15-stage column the slope criterion picks stage 10, the reflux handle picks stage 11, and the distillate handle picks stage 5 — a different section of the column, with a response fifteen times stronger. The same run shows the reflux response *changing sign* between stages 6 and 7, which no slope of the profile can reveal.

Choupo computes no relative-gain array, at any frequency: there is no steady-state gain matrix in the engine. The pairing heuristics below argue in terms of relative gains, and that argument is the literature’s, not a number this simulator produced.

Inventory first. Control every inventory (level, pressure) before any quality — an unbounded inventory shuts the plant down.

Controlled variable	Manipulate	Why
Level (drum, sump, tank)	the OUTflow	self-regulating; never control level by an inflow that sets production
Pressure (vapour)	vent / vapour outflow / condenser duty	fast; pins the vapour inventory
Feed flow / throughput	the feed valve directly	the master production handle
Reactor / column temperature	coolant or steam duty	the quality proxy
Column composition	a sensitive tray TEMPERATURE	direct composition is slow/expensive; tray T tracks it

Tuning starting points. Process PID is **mostly PI** — derivative amplifies noise, so add D only on slow, lag-dominated (temperature) loops. Flow loops: fast and tight. Level loops: loose/averaging (the tank is a buffer). Use the open-loop step (gain, dead time θ , time constant τ) with an IMC/ λ rule rather than trial and error.

Structure beyond single loops.

- **Cascade** when a fast inner disturbance exists (jacket-T inside reactor-T rejects coolant upsets early).
- **Feedforward** when a disturbance is MEASURED (feed rate/composition) — act before the error appears, trim with feedback.
- **Ratio** to hold two streams in proportion (stoichiometry, reflux, fuel/air).
- **Override / constraint** to ride a limit (max T, min level) safely.

Pairing discipline: pair each manipulated variable with the controlled variable it affects most directly and with least interaction (a near-1 relative gain). Two loops that fight oscillate — the glass-box dynamic run shows it at once. Keep DOF honest: one independent manipulated variable per controlled variable, no more.

Equipment-specific schemes. Beyond single loops, each unit has a conventional strategy.

Distillation is the hard one: a column has **five** handles (condenser duty, distillate D, reflux L, boilup V, bottoms B) and **five** things to hold (pressure, reflux-drum level, sump level, a top and a bottom composition). Pressure is fixed first (condenser duty or a vent). Each level **MUST** then be held by an OUTFLOW — and the rule is **control a level with the LARGER stream leaving that drum**: if the reflux ratio $R = L/D$ is high, reflux is the big stream, so hold reflux-drum level with reflux and set D for composition (the “LV” vs “DV” choice). Composition is inferred from a **sensitive tray TEMPERATURE** (fast, cheap) instead of a slow analyser. **Single-end** control (hold one composition, let the other float) is the robust default; **dual-composition** control saves energy but the



Figure 3: A conventional distillation control scheme: pressure (PC) on the condenser, each level (LC) held by an outflow, composition by a sensitive tray temperature (TC). Five handles, five controlled variables – the pairing is what makes it subtle.

two loops interact strongly (high relative gain) and usually need decoupling — which is exactly why students find column control hard. *Reactor*: temperature by coolant/steam duty, usually **cascade** (jacket-T inner, reactor-T outer). *Compressor*: an **anti-surge** recycle loop keeps flow above the surge line. *Fired heater*: outlet T trims the fuel, with an air/fuel **ratio** loop and a **low-flow trip** to protect the tubes.

20 Pressure relief, flare & safety basics

This is a Pareto orientation, not a safety course. The governing standards are **API STD 520/521/526** (sizing, scenarios, devices) and **API STD 537/NFPA 30**; treat the numbers below as where-to-start hypotheses, then run the controlling case.

When a device is mandatory. *Every* closed volume that can be isolated and over-pressured needs an independent overpressure-protection device — no exceptions for "small" vessels. The trigger is the credible-cause inventory, not size. A relief device is the last line of defence: it must work with zero operator action and zero instrument power, which is why a spring-loaded PSV (or rupture disc) is preferred over an instrumented trip for the primary case.

Set vs design pressure (the WHY). Design pressure is set ABOVE the maximum operating pressure so the PSV does not weep on normal swings. Heuristic: **design P = max(1.1 × operating P, operating P + 1.7 bar)**; for vessels under ~7 barg the +1.7 bar floor dominates, above it the +10% rule does. The **PSV set pressure = design pressure (MAWP)**. Code allows accumulation above set — typically **+10%** for a single non-fire device, **+16%** for multiple devices, and **+21%** for the fire case — because momentary overshoot during full relief is tolerable but permanent yield is not.

Sizing scenarios — size for the worst credible single cause (API 521).

Scenario	Typical controlling load	Why / note
External fire	Wetted-area heat input $\approx$ 43–70 kW/m² ($q = 43.2 \cdot F \cdot A^{0.82}$, SI)	Often controls vessels with large liquid inventory; sets a large orifice
Blocked outlet	Full pump/compressor/feed flow at shut-in	Common control case for pumped/compressed systems
Thermal expansion	Liquid trapped between two valves, heated	Tiny relief ($\frac{3}{4}$ "×1") — but mandatory; cheap to forget
Tube rupture (HX)	High-P side leaks to low-P side; full ΔP flow	The "2/3 rule": if low-side design P < 2/3 high-side, must relieve
Control-valve failure / loss of cooling	Failed-open bypass, reflux/condenser loss	Check fail-safe action of each valve

Size each, then **the largest required orifice wins**; pick the next standard API-526 letter orifice (D through T).

PSV vs rupture disc.

Device	Use when	Trade-off
Spring PSV	Reclosing wanted; clean service	Reseals, testable; can weep, can plug/foul
Rupture disc	Fast/large relief, fouling/corrosive, polymerising	Non-reclosing (full loss of inventory); cheap
Disc + PSV in series	Corrosive service needing reseal	Disc isolates/protects the PSV; best of both

Disposal — flare vs vent. Vent to atmosphere only if the relieved fluid is **non-toxic, non-flammable** at grade and the discharge disperses safely (e.g. steam, air, inert). Flammable/toxic streams route to a **closed flare header** → **knockout drum** → **flare**, sized so flare-tip exit velocity stays below **~0.5 Mach** and grade-level radiation stays under ****~6.3 kW/m²**** at the fence (API 521). Choupo models the relieved stream as a flash to header pressure; the **PSV-vent enthalpy and flare duty appear in the first-law ledger**, so the energy of blowdown is visible, not hidden.

Spacing/segregation (Pareto). Fired heaters and flares are ignition sources — keep **~15 m** from process units handling flammables; locate control rooms and the flare upwind/away. These setbacks are an insurance-grade default (IRI/GAP, API 521 §5), refined later by a dispersion/QRA study.

Governing standards — go to the source; lives depend on it. This section is an orientation only; the binding requirements live in, among others: pressure relief — **API STD 520 & 521**, tank venting — **API STD 2000**, overpressure protection — **ASME BPVC Section VIII** (UG-125+) and **Section I** (boilers); relief-device specs — **API STD 526**; flares — **API STD 521 & 537**; process-safety management — **OSHA 29 CFR 1910.119** and the **AIChE/CCPS** guideline series; functional safety / SIS — **IEC 61508 & 61511**; hazard-study method — **IEC 61882 (HAZOP)**; flammable liquids — **NFPA 30**; electrical area classification — **NFPA 70 / IEC 60079**;

facility siting — **API RP 752/753**. Size and select to the standard, then have it independently reviewed. See the Disclaimer & scope at the front of this guide.

21 Effluent treatment — liquid & gaseous

Every flowsheet has a back door: the streams you do NOT sell. Design them with the same care as the product train — environmental limits are *legal* limits, and an untreated vent or outfall can shut a plant down. **Hypothesis to run:** before sizing any treatment, MINIMISE the load at source (recycle, substitution, leak reduction) — the cheapest effluent is the one never made.

21.1 Liquid effluent (wastewater)

A train of increasing cost and specificity; stop when you meet the discharge limit (typically COD, BOD, suspended solids, pH, oil & grease, specific toxics):

Stage	Removes	Typical numbers	Use when
Equalisation + neutralisation	surges, pH	discharge pH 6–9; hours of buffer	always first
API separator / DAF	free oil, solids	oil to < 10–50 mg/L	oily / heavy-solids water
Biological (activated sludge)	BOD / COD	85–98% BOD removal; F/M 0.2–0.5	biodegradable organics
Adsorption (GAC) / advanced oxidation	refractory toxics	polishing to ppb	non-biodegradable, colour
Membrane (UF/RO) / evaporation	dissolved salts	water reuse, zero-liquid-discharge	reuse or brine concentrate

WHY a train: each stage is cheap *only* for what it targets — biology cannot take free oil, carbon is wasted on bulk BOD. Match the stage to the contaminant.

21.2 Gaseous effluent (air emissions)

Pollutant	Control	Typical numbers
Particulates	cyclone → bag filter / ESP	baghouse to < 10 mg/m ³ ; cyclone only > 5–10 μm
Acid gases (SO _x , HCl)	wet / dry scrubber	90–99% removal; L/G 5–15 L/m ³
VOC, dilute & high-flow	thermal / regenerative oxidiser	95–99% destruction at 760–820 °C; RTO recovers ~95% heat
VOC, concentrated / valuable	condensation / carbon / absorption	recover the solvent if it pays
NO _x	low-NO _x burner + SCR	SCR 80–90% at 300–400 °C with ammonia

The master rule: **dilute-and-high-flow favours destruction (oxidiser); concentrated-and-valuable favours recovery (condense/absorb)** — the same see-then-decide trade-off as everywhere else. Choupo's cyclone, absorber, and combustion duties size the front end; the discharge limit is the spec and the local regulation is the law (see *Standards & further reading*).

22 Cost estimation & profitability

Costing in Choupo is a *post-processing pass* (**CostingPass**, Guthrie/Turton models) feeding **EconomicsPass** (see **optim05_reactor_npv**). The numbers below are heuristics to be *run and checked*, never final verdicts — accuracy depends entirely on the AACE estimate class.

Capital (CAPEX). Two routes, increasing in effort:

- *Lang factor* — fixed capital = Lang factor $\times \Sigma$ (delivered equipment cost). Hypothesis to verify: a fluids plant lands near $4.7\times$; solids-handling near $3.1\times$; mixed $3.6\times$. The factor absorbs piping, instrumentation, civil, and indirects, so it is crude (± 30 -40%) but fast.
- *Bare-module (Guthrie/Turton)* — per-unit: $C_{BM} = C_p^0 F_{BM}$, with $F_{BM} = (B_1 + B_2 F_M F_P)$. Pressure (F_P) and material (F_M) multipliers are *itemised*, so this is the glass-box default. Verify: a SS316 vessel at 30 bar shows $F_{BM} \approx 5 - 7$ vs ~ 3 for carbon steel at 1 bar.

Item	Typical range	Why
Lang factor (fluids)	4.5-5.0	Indirects dominate fluid plants
Lang factor (solids)	3.0-3.6	Less piping/instrumentation
Bare-module F_{BM} (CS, low-P)	2.5-4.0	Base installation overhead
Six-tenths exponent n	0.5-0.7 (≈ 0.6)	Economy of scale on area/volume
CEPCI escalation	index ratio	~ 600 (2019) $\rightarrow \sim 800$ (2023)
Site/auxiliary uplift	+10-25%	Grassroots vs battery-limit

Six-tenths rule: $C_2 = C_1(S_2/S_1)^{0.6}$. Doubling capacity raises cost only $\sim 52\%$ ($2^{0.6}$) — verify before sizing a redundant train. *CEPCI:* $C_{now} = C_{base}(I_{now}/I_{base})$; skipping a 4-year escalation understates CAPEX $\sim 25\%$.

Operating cost (COM). Rule of thumb (Turton): $COM \approx 0.18 FCI + 2.73 C_{OL} + 1.23(C_{UT} + C_{WT} + C_{RM})$. Verify the dominant term:

COM component	Typical share	Why
Raw materials	40-80%	Usually dominant; selectivity is king
Utilities	10-30%	Steam/cooling/power — pinch can cut 20-30%
Labour + supervision	5-15%	Headcount $\times$ shift factor 4.5
Maintenance	3-6% of FCI	Wear-driven

Because raw materials dominate, a 1% yield gain often beats any utility optimisation — test it against a pinch retrofit in the same case.

Profitability. Three lenses, never one:

- *Payback period* (FCI / annual after-tax cash flow): attractive < 3 yr, marginal 3-5 yr, suspect > 5 -6 yr. Ignores time value, so it is a screen only.
- *NPV* at a hurdle rate (10-15% typical): positive $\Rightarrow$ value-adding; verify sensitivity to discount rate and a 10-year horizon.
- *IRR/DCFROR*: the rate where NPV=0; want it 5-10 points above WACC. Below the hurdle even a short payback fails.

Estimate class governs everything. Report the band, not a point:

AACE class	Maturity	Accuracy	Method
Class 5	0-2%	$-30/+50\%$	Capacity-factored, six-tenths
Class 4	1-15%	$-20/+30\%$	Equipment-factored (Lang)
Class 3	10-40%	$-15/+20\%$	Bare-module / budget
Class 2	30-70%	$-10/+15\%$	Detailed, semi-quantitative

AACE class	Maturity	Accuracy	Method
Class 1	50-100%	-5/+10%	Definitive, quote-based

Quoting an NPV to four significant figures off a Class-5 estimate is the classic glass-box failure: the $\pm 50\%$ CAPEX band swamps the answer. Always run the case at both error bounds.

23 Materials of construction & mechanical design

Material choice is a *climb-the-ladder* decision: start with the cheapest metal that survives the duty, and step up only when corrosion rate, chlorides, temperature, or hydrogen force you. Each rung is roughly $1.5\text{--}4\times$ the cost of the one below, so an unjustified jump from carbon steel to SS316 can double a vessel's bare cost (see Choupo's Guthrie/Turton material-factor f_M in the costing pass; carbonSteel/SS304/SS316/aluminium are the reference data/materials/*.dat).

The selection ladder. The governing rule of thumb is a target corrosion rate $< 0.1 \text{ mm/yr}$ ($\approx 5 \text{ mil/yr}$) for a 20-year life; $0.1\text{--}0.5 \text{ mm/yr}$ is tolerable with a fat allowance; above 0.5 mm/yr step up.

Service / limit	Pick	Why (the WHY)
Dry hydrocarbons, steam, mild T < 400 °C	Carbon steel	Cheapest; $\sim 0.1 \text{ mm/yr}$ in benign duty
Dilute acids, oxidising, T < 60 °C, chloride-free	SS304	Cr_2O_3 passive film; loses it to Cl^-
Chlorides, dilute H_2SO_4 , food/pharma	SS316 (2–3% Mo)	Mo resists pitting; PREN ≈ 24 vs 18
Hot chloride + tensile stress	Duplex 2205 / Ni alloy	Austenitic SCC kills 304/316
Strong HCl, wet Cl_2 , HF	Hastelloy C-276 / lined	General attack defeats stainless
High-purity water, weight-critical	Aluminium	Cheap, light, but soft & low-T only

Chlorides & stress-corrosion (the silent killer). Austenitic 304/316 suffer chloride **SCC** above $\approx 60 \text{ °C}$ with tensile stress and even ppm-level Cl^- — failures look like brittle cracking with no wall loss, so a corrosion allowance does *not* protect you. Below 60 °C or for hot chlorides, jump to duplex (ferritic-austenitic, immune) or nickel alloys. Pitting resistance scales with **PREN** = $\% \text{Cr} + 3.3 \cdot \% \text{Mo} + 16 \cdot \% \text{N}$ (304 ≈ 18 , 316 ≈ 24 , 2205 ≈ 35).

High-T creep & hydrogen. Carbon steel is creep-limited to $\approx 425 \text{ °C}$; above this design stress collapses, so $1\frac{1}{4}\text{Cr}\text{--}1\frac{1}{2}\text{Mo}$, 9Cr, or austenitics take over (relevant for fired-heater tubes and HiTec-salt service to 535 °C). In H_2 at pressure, consult **Nelson curves** (API 941): below the curve carbon steel is safe; above it, hydrogen attacks carbides $\rightarrow$ decarburisation and fissuring, demanding Cr-Mo steels.

Corrosion allowance. Add **3.0 mm** standard (1.5 mm benign, up to 6 mm aggressive). It is sacrificial metal, *not* an SCC or hydrogen fix.

Wall thickness vs pressure. Thin-wall (ASME VIII Div. 1): $t = P \cdot R / (S \cdot E - 0.6 \cdot P) + CA$, with weld efficiency $E = 0.85\text{--}1.0$ and S the allowable stress ($\approx 138 \text{ MPa}$ for SS316 at moderate T). Design pressure margin: **max(operating + 10%, +1.7 bar)**; design temperature: operating + **25 °C**. A minimum shell of $\sim 6\text{--}8 \text{ mm}$ is set by handling, not pressure.

Aspect ratios. Vertical process vessels and knockout drums: $L/D \approx 3$ (range 2–5) — minimises wall metal at fixed volume while giving disengaging height. Distillation columns run L/D up to **20–30**, capped by wind/seismic moment. Atmospheric storage tanks favour $L/D \approx 0.5\text{--}1$ (squat, cheap floor). Horizontal drums sit at $L/D \approx 3\text{--}4$.

Glass-box check: re-run the Choupo sizing/costing pass swapping material carbonSteel; $\rightarrow$ SS316; and watch f_M and shell mass move; that is the cost of one rung.

Corrosion allowance — how many millimetres? Wall thickness = the pressure-stress thickness (from ASME) **plus a sacrificial corrosion allowance**: $CA = \text{corrosion rate} \times \text{design life}$. Carbon steel corroding at $\sim 0.1\text{--}0.25 \text{ mm/yr}$ over a 20–25 year life gives the familiar **3 mm** (1/8 in) allowance; severe service runs to **6 mm**, clean/inert service to **0–1.5 mm**, and a corrosion-resistant alloy needs **none** (it is chosen so the rate is ~ 0).

Corrosion rate	Verdict
< 0.1 mm/yr	excellent — carbon steel, 1.5–3 mm allowance
0.1–0.5 mm/yr	acceptable with allowance + inspection
> 0.5 mm/yr	do NOT just thicken the steel — upgrade the alloy or line it

The rule: a corrosion allowance buys time, not safety — above $\sim 0.5 \text{ mm/yr}$ the economics flip to a better material.

Gaskets — the small part that floods the floor. A gasket seals a flange by being squeezed between the faces; it must survive the **pressure, the temperature, AND every fluid the joint ever sees — including cleaning**. Select by P–T–fluid:

Gasket	P–T window	Use when
Compressed fibre / PTFE sheet	to ~20 bar, ~200 °C	low duty; PTFE for chemical resistance
Spiral-wound (metal + graphite/PTFE filler)	to ~100+ bar, ~450 °C+	the workhorse for process flanges
Ring-joint (RTJ, solid metal)	very high P (> 100 bar)	high-pressure gas, wellheads

The wetted material — the winding and the filler — must resist the worst fluid, and the worst fluid is often **not** the product but the **CLEANING** step. A gasket chosen only for cold wine will corrode through during a **40 °C acid CIP wash**, and then the wine runs out of the joint on the next batch — a real and classic failure. **Specify the gasket (and the bolt load) for the most aggressive condition the joint ever sees, cleaning included**, never for normal operation alone.

24 Standards & further reading

This guide distils the classic references below. For any real design go to the primary source — and for anything safety- or environment-critical, the governing standard is the authority, not this guide (see *Disclaimer & scope*).

General process design & heuristics

- Turton, Shaeiwitz, Bhattacharyya & Whiting, *Analysis, Synthesis, and Design of Chemical Processes*, Prentice Hall.
- Seider, Seader, Lewin & Widagdo, *Product and Process Design Principles*, Wiley.
- Sinnott & Towler, *Chemical Engineering Design* (Coulson & Richardson, Vol. 6), Butterworth-Heinemann.
- Green & Perry, *Perry's Chemical Engineers' Handbook*, McGraw-Hill.
- Douglas, *Conceptual Design of Chemical Processes*, McGraw-Hill.
- Branan, *Rules of Thumb for Chemical Engineers*; Walas, *Chemical Process Equipment*.

By topic

- Heat integration / pinch — Smith, *Chemical Process Design and Integration*, Wiley; Kemp, *Pinch Analysis and Process Integration*.
- Distillation — Kister, *Distillation Design & Distillation Operation*, McGraw-Hill.
- Heat transfer — Kern, *Process Heat Transfer*; **TEMA** standards.
- Reactors — Fogler, *Elements of Chemical Reaction Engineering*; Levenspiel, *Chemical Reaction Engineering*.
- Membranes — Baker, *Membrane Technology and Applications*, Wiley.
- Materials / corrosion — **NACE/AMPP** standards; **API RP 571** (damage mechanisms).
- Effluent — Metcalf & Eddy, *Wastewater Engineering*; **US EPA** and **EU Industrial Emissions Directive (BAT / BREF)** guidance.

Safety & environmental standards — binding, not optional

- Pressure relief & venting — **API STD 520 / 521 / 526 / 2000**; **ASME BPVC Sec. VIII & I**.
- Flares — **API STD 521 / 537**.
- Process-safety management — **OSHA 29 CFR 1910.119**; **AICHE CCPS** guideline series.
- Functional safety / SIS — **IEC 61508 / 61511**; **HAZOP — IEC 61882**.
- Flammables — **NFPA 30**; electrical area classification — **NFPA 70 / IEC 60079**.
- Emissions & discharge — the local regulator's limits (e.g. **US EPA**, **EU IED / BAT-BREF**).

25 Equipment selection at a glance

26 Equipment selection at a glance

This one-page table is the *index of bets* — each row is a duty you will face, the default pick, the one-line reason, and the range to start from. Treat every "typical number" as a hypothesis to verify by running the linked Choupo case. When two picks compete, the table gives the **default**; the section explains the crossover.

Duty / situation	Pick this	Because	Typical numbers
Overall flowsheet structure	Reactor → separation → recycle, with a purge on any loop that accumulates inerts	The "onion": fix conversion + recycle before heat/utilities; purge bleeds what recycle would otherwise trap	Purge loses 1–5 % of recycle; target recycle so single-pass conversion $\neq$ overall conversion
Equilibrium-limited reaction	Push reactor to ~90–95 % of equilibrium, recycle the rest	Last few % of approach costs disproportionate volume; recycle is cheaper than oversizing	Single-pass $X = 0.6\text{--}0.9$; overall $X > 0.99$ via recycle
Fast liquid-phase reaction, high conversion	PFR (or CSTR-in-series)	Plug flow needs least volume for a given conversion; back-mixing penalises high X	PFR vol $\approx 1/3\text{--}1/2$ of one CSTR at $X = 0.9$
Autocatalytic, biological, solids-laden, or strong heat-removal duty	CSTR (or CSTR train)	Back-mixing dilutes feed, eases T control and solids handling	3–4 CSTRs in series $\approx$ PFR; τ set by k and target X
Gas-phase catalytic reaction	Packed-bed (PFR model)	Heterogeneous catalyst + plug flow; watch ΔP across bed	Superficial gas 0.3–1 m/s; bed ΔP a few % of P
Separate close-boiling liquids	Distillation	Cheapest workhorse whenever relative volatility allows	Viable for $\alpha > \sim 1.05$; easy for $\alpha > 1.5$
Distillation column operating point	Reflux $R = 1.1\text{--}1.3 \times R_{\min}$	Optimum trades reflux (energy) vs trays (capital); flat-bottomed cost curve	$R/R_{\min} \approx 1.1\text{--}1.3$; $N \approx 2\text{--}2.5 \times N_{\min}$
Heat-sensitive / high-boiling bottoms	Vacuum distillation	Lowers boiling T below decomposition; protects product	Run at 0.05–0.3 bar when $T_{\text{bottom}} > \sim 150\text{--}200$ °C at 1 atm
Column vapour loading	Size to F -factor / fraction of flooding	Avoids flooding (carry-over) and weeping	$F = u \cdot \sqrt{\rho_v} \approx 1\text{--}2.5$; design at 70–85 % of flood
Dilute aqueous dewatering / desalination	Reverse osmosis (spiral-wound)	Pressure-driven, no phase change; rejects monovalent salts	ΔP 15–70 bar (seawater); flux 15–30 LMH
Recover solvent / reject divalent ions, pass monovalent	Nanofiltration	Looser membrane, lower pressure than RO, partial-rejection selectivity	ΔP 5–20 bar; rejects multivalent > 90 %, monovalent 20–70 %
Remove particles from a gas	Cyclone (coarse) → bag filter (fine)	Cyclone is cheap above $\sim 5\text{--}10$ μm ; fabric filter takes sub-micron	Cyclone η high $> 5\text{--}10$ μm ; bag filter to < 1 μm
Produce a crystalline solid	MSMPR cooling crystalliser (or evaporative if solubility flat with T)	Cooling exploits steep solubility- T slope; evaporative when it is flat	Supersaturation low (metastable); residence min–hours
Clean liquid–liquid heat transfer	Shell-and-tube (or gasketed plate if T, P moderate)	Robust, costed, high U ; plates beat it on area/cost when clean & $< \sim 25$ bar	$U \approx 250\text{--}750$ W/m ² K; $\Delta T_{\min} \approx 10$ °C
Heat transfer with a gas on one/both sides	Extended/compact surface — not a bare shell-and-tube	Gas-side h is the bottleneck; gas–gas shell-and-tube needs huge area	$U_{\text{gas–gas}} \approx 10\text{--}50$ W/m ² K; $\Delta T_{\min} \approx 20$ °C

Duty / situation	Pick this	Because	Typical numbers
Condense overhead / reboil bottoms	Kettle/thermosiphon reboiler; shell-side condenser	Phase change gives very high h ; match utility to T-level	$U_{\text{cond}} \approx 1000\text{--}3000 \text{ W/m}^2\text{K}$; $\Delta T_{\text{min}} \approx 5\text{--}10 \text{ }^\circ\text{C}$
Process duty above $\sim 400 \text{ }^\circ\text{C}$	Fired heater (very high T) or molten-salt loop (250–535 $^\circ\text{C}$)	Steam tops out $\sim 250 \text{ }^\circ\text{C}$; salt covers the gap, fired heater above	HiTec salt 142–535 $^\circ\text{C}$; Dowtherm A to $\sim 400 \text{ }^\circ\text{C}$
Size a pumped liquid line	Velocity 1–3 m/s (suction lower)	Balances pumping ΔP vs pipe capital; avoids settling/erosion	Discharge 1–3 m/s; suction 0.5–1.5 m/s; $\Delta P \sim 0.5 \text{ bar/100 m}$
Size a gas/vapour line	Velocity 15–30 m/s, capped by ρu^2	Low-density fluid $\rightarrow$ high velocity OK; erosion-corrosion limits ρu^2	15–30 m/s; steam 30–60 m/s; $\rho u^2 < \sim 10000 \text{ kg/m} \cdot \text{s}^2$
Default material of construction	Carbon steel	Cheapest; specify up only when corrosion/T demands it	Use to $\sim 425 \text{ }^\circ\text{C}$, non-corrosive, no chlorides
Chloride or aqueous-corrosive service	SS316 (not SS304)	Mo content resists chloride pitting that defeats 304	316 for chlorides/acids; 304 for mild aqueous/atmospheric

Use this guide as a map, not a destination. Every range here is a place to *start* a Choupo case, not a number to copy onto a P&ID. The discipline the guide is teaching is the loop itself: form the heuristic bet, build the smallest case that tests it, run it, read what the solver actually does, and let the KPIs confirm or move the range for your system. The rules that survive that loop become judgment; the ones that fail teach you where your problem lives outside the textbook's assumptions. That is the whole point of a glass-box simulator — and the reason a heuristic you have personally falsified once is worth ten you merely memorised. Read the sections in the onion order, keep the at-a-glance table beside the keyboard, and verify everything by running it.